

recon

User Manual

Version 0.101.1

2026-06-18

Repository · <https://github.com/codedeviate/recon>

License · MIT

About this manual

recon is a versatile network reconnaissance CLI written in Rust. It started as a curl clone and grew into a multi-protocol investigation tool covering HTTP(S), TLS certificate inspection, DNS, WHOIS, ping, traceroute, barcode encoding and decoding, file compression and archiving, markdown/HTML/PDF conversion, and a full Rhai script engine that exposes every protocol probe and helper for automation.

This manual covers every user-facing flag, every script binding, and shows how to combine them. The built-in `recon --help <topic>` command remains the quick reference; this document is the long-form companion.

A PDF rendering of this manual lives alongside it at [MANUAL.pdf](#). It is regenerated every time the markdown changes — see `CLAUDE.md` for the maintenance policy. Generate locally with:

```
recon --md-to-pdf docs/MANUAL.md \  
  --toc --toc-depth 3 --gfm --unsafe-html --page-break-on-h1 \  
  --doc-title 'recon User Manual' \  
  --doc-author 'Thomas Bjork' \  
  --doc-subject 'recon CLI reference and script-engine guide' \  
  --doc-keywords 'recon, curl, network, reconnaissance, scripting' \  
  -o docs/MANUAL.pdf
```

Table of Contents

Contents

[About this manual](#)

[Table of Contents](#)

[Introduction](#)

[Installation](#)

[From source](#)

[First-run bootstrap](#)

[Quick start](#)

[How recon is organized](#)

[Global conventions](#)

[Part II — CLI reference](#)

[HTTP / HTTPS requests](#)

[Core flags](#)

[Examples](#)

[Wget-style batch fetching](#)

[Examples](#)

[Output control](#)

[Examples](#)

[Authentication & TLS](#)

[Examples](#)

[Client certificates \(mTLS\)](#)

[Examples](#)

[Browser fingerprint impersonation \(0.77.0, opt-in\)](#)

[Examples](#)

[Proxy routing](#)

[Examples](#)

[Unix-domain sockets](#)

[Examples](#)

[HSTS persistent cache](#)

[Examples](#)

[IPFS / IPNS](#)

[Examples](#)

[Cookies](#)

[Examples](#)

[DNS](#)

Examples

TLS certificate inspection

Examples

Configuration files

Environment overrides

CLI flags

Deep-merge rules

Worked example

Sections

Aliases / compatibility modes

Ping, traceroute, netstatus

Examples

Email protection

Examples

SMTP

Examples

Mail retrieval (POP3, IMAP)

Examples

File transfer

FTP flags

TFTP flags

Examples

Other protocols

SSH, SCP

Telnet

LDAP

WebSocket

RTSP

Dict (RFC 2229)

NTP, Memcached, Redis, MQTT

Source comparison

Examples

Document conversions

Examples

Cover pages

Page breaks

PDF page export

Barcode encoding & decoding

Examples

--encode-hints — Aztec / PDF417 tuning

HRT (human-readable text under 1D barcodes)

- Hashing
 - Examples
- Compression & archiving
 - Examples
- Encryption
 - Examples
- Check digits
 - Examples
- Sample data
 - Examples
- JWT tokens
 - Examples
- Text encoding
 - Examples
- Serve mode
 - Examples
- Write-out format
 - Examples
- Browser automation
 - Examples
- Meta flags
 - flags — the quick lookup
- Interactive REPL
 - Launching
 - Meta-commands
 - Multi-line input
 - Autoprint
 - Threading caveat
 - Example session
- Part III — Script engine
 - Running scripts
 - Bare-word invocation
 - Script language (Rhai)
 - Types and literals
 - Operators
 - Control flow
 - Functions and closures
 - Error handling
 - Standard built-in functions
 - Shebang — executable scripts
 - CLI inheritance

Core helpers

Examples

Output bindings

Examples

File I/O bindings

Whole-file helpers

Streaming handle API

Examples

Clipboard bindings

Examples

Comparison bindings

Examples

HTTP binding

Response map

Options map

Examples

Browser (sticky-session) binding

Examples

TCP, TCP-server, UDP

TCP client (existing)

TCP server (0.57.0)

UDP

Examples

Threading bindings

Examples

Shell binding

Return value (shell)

Opts map (both forms)

Examples

TUI dashboard binding

Example

v1 limitations (intentional)

DNS binding

Examples

TLS probe binding

Examples

Ping binding

Examples

File transfer bindings

FTP

SFTP

TFTP

Gopher

Examples

Mail retrieval bindings

POP3

IMAP

Examples

SMTP binding

Examples

WebSocket binding

Examples

Other protocol bindings

LDAP

RTSP

Dict

NTP

Memcached

Redis

MQTT

Encoding & decoding bindings

Examples

Hashing binding

Examples

Compression & archive bindings

Compression

Archive

Examples

Encryption bindings

Examples

Check-digit binding

Examples

Sample-data binding

Examples

JWT binding

Examples

Email-protection binding

Examples

Netstatus binding

Text-encoding binding

Examples

String helpers

Examples
jq filter
Example
git wrapper
Example
gh wrapper
Auto-account-switch
Methods
Example
Whois binding
IPFS binding
Agent-browser binding
Existence check
Examples
Global options (0.75.0)
SQLite binding
Examples
Document-conversion bindings
Opts map
Examples
pdf_export_page
AI bindings (ai::*)
Backend selection
Config file
Example
Errors
Telemetry (-v / -vv)
Part IV — Appendices
Exit codes
Environment variables
Configuration file
Structure
~/.recon/ layout
Glossary

The table of contents above is auto-generated from H1/H2/H3 headings. Every entry is a clickable anchor in the PDF. For a narrative walkthrough of the structure, the four parts are:

- **Part I — Getting started:** introduction, installation, quick start.
- **Part II — CLI reference:** every flag grouped by area, with examples.
- **Part III — Script engine:** the Rhai interpreter, every binding, many examples.
- **Part IV — Appendices:** exit codes, env vars, config, glossary.

Introduction

recon is a single-binary command-line tool. Its design goals, in rough order:

1. **Curl-compatible** where it overlaps with curl. Flag names, short forms, and behaviours match curl for HTTP(S) requests so existing tooling keeps working. `recon --help curl` lists the compatibility status.
 2. **Pure Rust** where possible. request + rustls for HTTP(S); hickory for DNS; comrak for markdown; flate2/brotli/zstd for compression; rxing for barcodes; tungstenite for WebSockets; rumqttn for MQTT. Where a pure-Rust path doesn't exist (headless-Chrome PDF rendering, for instance), recon shells out to a named external CLI (`agent-browser`) rather than linking the external dependency.
 3. **Protocol-rich**. HTTP(S), FTP(S), SFTP, TFTP, Gopher, SMTP(S), POP3(S), IMAP(S), SSH, SCP, Telnet, LDAP(S), WS(S), RTSP(S), Dict, NTP, Memcached, Redis, MQTT(S), IPFS/IPNS, Unix sockets, plus ping/traceroute.
 4. **Scriptable**. A Rhai script engine exposes every protocol probe, plus cryptography, compression, barcodes, JSON/YAML/XML handling, SQLite, file I/O, and multi-threaded concurrency. Scripts can stand in for shell pipelines and integration-test harnesses.
 5. **Diagnostic-friendly**. Verbose output, curl-style write-out, prettified JSON/XML, precise error messages, exit codes that match curl's.
-

Installation

From source

```
git clone https://github.com/codedeviate/recon
cd recon
cargo build --release
cp target/release/recon ~/.local/bin/      # or /usr/local/bin with sudo
```

You'll need:

- Rust 1.75 or newer (`rustup install stable`).
- A system linker; `cargo` handles the rest.
- For the optional `--html-to-pdf` / `--md-to-pdf` features: `agent-browser` on `$PATH` (`brew install agent-browser` on macOS, or `npm install -g agent-browser` elsewhere).

First-run bootstrap

```
recon --init
```

creates `~/.recon/` with `script/`, `jars/`, `sni/`, and a commented `config.toml` skeleton. Existing files are never overwritten. See the [~/.recon/ layout](#) appendix.

Quick start

```

# HTTP(S)
recon https://api.example.com/status
recon -X POST https://api.example.com/users -d '{"name":"a"}' -H 'Content-Type: application/json'
recon --json '{"q":"rust"}' https://api.example.com/search # auto-sets headers
recon -L https://short.url/foo # follow redirects
recon https://site.example.com --LHEAD # follow
+ show each hop

# Inspection
recon https://example.com --cert # inspect TLS cert
recon example.com --dns #
A/AAAA/MX/TXT/NS/CNAME/SOA
recon --ping 8.8.8.8 # TCP + ICMP ping
recon --traceroute 8.8.8.8
recon --whois example.com

# Email-protection
recon example.com --spf --dmarc --dkim selector1 --mta-sts --bimi --tls-rpt

# Barcodes
recon --encode qr -o out.png 'https://example.com'
recon --decode out.png # →
qr<TAB>https://example.com

# Conversions
recon --md-to-pdf README.md --toc --gfm -o README.pdf
recon --compare a.json b.json # GNU-diff compatible

# Scripting
recon --script my-probe.rhai https://api.example.com
recon --init #
bootstrap ~/.recon/

# Help surface
recon --help # clap-generated flag list
recon --flags # curl-style alphabetical index
recon --help http # deep-dive on HTTP topic
recon --help script # deep-dive on scripting

```

```
recon --examples # curated
examples, paged
recon --version #
protocols + features
```

How recon is organized

recon has one binary with many modes. The mode is selected by flag:

Mode group	Selector flags
HTTP request (default)	no mode flag, or a URL positional argument
Source-inspection	<code>--hash</code> , <code>--compress</code> , <code>--decompress</code> , <code>--encode</code> , <code>--decode</code> , <code>--encrypt</code> , <code>--decrypt</code>
Email-protection	<code>--spf</code> , <code>--dmarc</code> , <code>--dkim</code> , <code>--mta-sts</code> , <code>--bimi</code> , <code>--tls-rpt</code>
Network tests	<code>--ping</code> , <code>--traceroute</code> , <code>--netstatus</code> , <code>--whois</code>
Certificate inspection	<code>--cert</code> , <code>--dns</code> , or positional URL with <code>--cert</code> / <code>--dns</code>
SMTP send	<code>--mail-from</code> and <code>--mail-to</code>
Serve	<code>--serve</code> , <code>--serve-tls</code>
JWT	<code>--jwt-view</code> , <code>--jwt-sign</code> , <code>--jwt-validate</code>
Archive	<code>--archive</code> , <code>--extract</code>
Checkdigit	<code>--checkdigit</code> , <code>--checkdigit-create</code> , <code>--checkdigit-list</code>
Sample	<code>--sample</code> , <code>--sample-list</code>
Browser	<code>--browser-screenshot</code>
Compare	<code>--compare</code>
Docs	<code>--md-to-html</code> , <code>--md-to-pdf</code> , <code>--html-to-pdf</code>
Script	<code>--script</code> (turns recon into a Rhai interpreter)
Meta	<code>--init</code> , <code>--editor-cleanup</code> , <code>--list-charsets</code> , <code>--iconv</code>

Only the modes in the first group do an HTTP request. Every other mode is stand-alone.

Global conventions

- **-h, --help** prints the clap-generated flag summary. **--help <topic>** routes to a deep-dive help topic (see `recon --help` footer for the list).
 - **-V, --version** prints the curl-compatible multi-line banner. **--version-short** prints just the version number.
 - **-v** increases verbosity. Repeat (**-vv** , **-vvv**) for more detail. At default, connection lines and protocol handshakes are suppressed.
 - **-s, --silent** suppresses progress + verbose output. **-S, --show-error** re-enables error output when **-s** is set (curl-compat). Use **--status** (long form only) to print just the HTTP status code.
 - **-o <FILE>** writes the response body to a file instead of stdout. **-O / --remote-name** derives the filename from the URL's path. **-J / --remote-header-name** respects Content-Disposition.
 - **-f, --fail** exits non-zero on HTTP 4xx/5xx (no body printed). **--fail-with-body** prints the body but still exits non-zero.
 - **-k, --insecure** disables TLS cert verification. Use for self-signed tests; never in production scripts.
 - **--pretty** pretty-prints JSON / XML / YAML bodies. Auto-detection by Content-Type; can be disabled with **--no-pretty**. Long form only — **-p** is reserved for curl's **--proxytunnel**.
 - **-L, --location** follows HTTP redirects. **--max-redirs N** caps the chain (default 10). **--LHEAD** follows and prints each hop's headers.
 - **-A "..."** sets the User-Agent. Default is `recon/<version>`.
 - **-e <URL>** sets the Referer header.
 - **-H 'Name: Value'** adds a request header. Repeat for multiple.
 - **--max-time <SECS>** overall operation timeout. **--connect-timeout <SECS>** TCP connect timeout (default 30).
 - **-w '<FORMAT>'** writes a curl-compatible write-out string to stdout after the body. See [Write-out format](#).
 - **--json <DATA>** curl's **--json** compatibility: sends DATA as a JSON body and auto-sets `Content-Type: application/json + Accept: application/json`.
 - **-T <FILE>** uploads the file as the request body (default method PUT).
 - **Environment variables** — many flags honor the same env var curl would (`HTTP_PROXY` , `HTTPS_PROXY` , `NO_PROXY` , `ALL_PROXY` , `SSL_CERT_FILE` , `CURL_CA_BUNDLE`). recon-specific ones are prefixed `RECON_`. See [Environment variables](#).
-

Part II — CLI reference

HTTP / HTTPS requests

The default mode. A positional URL (or `--url <URL>`) triggers an HTTP request through the request + rustls stack.

0.62.0 curl-parity additions — see the curl easy-wins section of `recon --examples` and the relevant flags table rows below. Quick summary: `-r/--range`, `-z/--time-cond`, `--etag-compare`, `--etag-save`, `--timestamping`, `--max-filesize`, `--url-query`, `--disallow-username-in-url`, `--remove-on-error`, `--no-clobber`, `--create-file-mode`, `-N/--no-buffer`, `-D/--dump-header`, `--stderr`, `--no-progress-meter`, `--tcp-nodelay`, `--no-keepalive`, `--keepalive-time`, `--capath`, `--ca-native`, `--tls-max`, `--connect-to`, `--oauth2-bearer`, `--xattr`, `--spider`.

Core flags

Flag	Description
<code>-X, --request <METHOD></code>	GET, POST, PUT, PATCH, DELETE, HEAD. Default GET (PUT when <code>-T</code> set, POST when <code>-d</code> set).
<code>-H, --header <NAME: VALUE></code>	Add a request header. Repeatable.
<code>-d, --data <BODY @FILE></code>	Request body. <code>@file</code> reads from disk, <code>@-</code> reads stdin. Promotes GET → POST unless <code>-G</code> .
<code>--json <DATA></code>	Curl-compatible: sends DATA as JSON, sets <code>Content-Type: application/json + Accept</code> .
<code>--data-raw <DATA></code>	Like <code>-d</code> but <code>@file</code> is NOT processed (sends literal string).
<code>--data-binary <DATA></code>	Like <code>-d</code> but CR/LF are not stripped from <code>@file</code> content.
<code>--data-urlencode <DATA></code>	URL-encode DATA. Repeatable; values joined with <code>&</code> .
<code>-G, --get</code>	Send <code>-d</code> data as query parameters on GET (body empty).
<code>-T, --upload-file <PATH></code>	Upload file as request body. Default PUT.
<code>-L, --location</code>	Follow redirects (3xx).

Flag	Description
<code>--max-redirs <N></code>	Cap redirect chain (default 10).
<code>--LHEAD</code>	Follow redirects and print each hop's headers.
<code>-e, --referer <URL></code>	Set Referer header. <code>--referrer</code> alias accepted.
<code>-A, --user-agent <STR></code>	User-Agent. Default <code>recon/<version></code> .
<code>--compressed</code>	Request + auto-decompress gzip/deflate/brotli/zstd.
<code>--max-time <SECS></code>	Total operation timeout (seconds, fractional OK).
<code>--connect-timeout <SECS></code>	TCP connect timeout (default 30).
<code>-f, --fail</code>	Exit non-zero on HTTP 4xx/5xx, suppress body.
<code>--fail-with-body</code>	Exit non-zero on 4xx/5xx but still print body.
<code>-4, --ipv4 / -6, --ipv6</code>	Force IPv4 / IPv6 resolution.
<code>--resolve <HOST:PORT:ADDR></code>	Override DNS for a specific host:port → addr. Repeatable.

Examples

```

# Simple requests
recon https://httpbin.org/get
recon -X POST https://api.example.com/v1/users -d '{"name":"a"}' -H
'Content-Type: application/json'
recon --json '{"query":"rust"}' https://api.example.com/search

# Body from files / stdin
recon https://upload.example.com -d @payload.json -H 'Content-Type:
application/json'
cat body.json | recon -X POST https://api.example.com -d @-

# Upload
recon -T ./report.pdf https://upload.example.com/reports/          # PUT by
default
recon -T ./report.pdf https://upload.example.com/ -X POST          # override
method

# Follow redirects
recon -L https://bitly.example.com/abc
recon --LHEAD https://example.com/                                # show each
hop's headers

# Fail fast
recon -f https://api.example.com/users/42                          # exit 22
on 4xx/5xx, no body
recon --fail-with-body https://api.example.com/users/42           # same but
still print body

# Rate and timeout
recon --max-time 10 https://slow.example.com/
recon --connect-timeout 3 --max-time 8 https://flaky.example.com/
recon --limit-rate 500K https://download.example.com/big.bin -o big.bin
recon --speed-limit 10000 --speed-time 30
https://stalling.example.com/big.bin -0

```

Wget-style batch fetching

Long-form wget-compatible flags for batch URL handling. Short forms (`-A`, `-R`, `-w`, `-t`, `-r`, `-l`, `-m`, `-p`, `-k`, `-D`, `-H`, `-np`) are intentionally not provided — recon reserves single-letter flags for curl compatibility. The recursive/mirror cluster (`-r`, `-l`, `-m`, `-p`, `-k`, etc.) is deferred and tracked in `OUT-OF-SCOPE.md`.

Flag	Description
<code>--input-file</code> <code><FILE></code>	Batch-fetch URLs listed in FILE (one per line, <code>#</code> comments, blank lines ignored, <code>-</code> reads from stdin). Each URL processed

Flag	Description
	independently.
<code>--wait <SECS></code>	(0.67.0) Fixed-seconds delay between URLs in batch mode. Skipped before the first URL. Overrides <code>--rate</code> when both are set.
<code>--tries <N></code>	(0.67.0) Total attempts per URL (wget semantics: <code>tries = retries + 1</code>). Overrides <code>--retry</code> . <code>--tries 1</code> disables retries; <code>--tries 0</code> is rejected.
<code>--accept <LIST></code>	(0.67.0) Comma-separated filename-suffix accept list (case-insensitive). e.g. <code>--accept jpg,png</code> keeps only URLs whose final path segment ends in <code>.jpg</code> or <code>.png</code> . URLs with empty final segments fail.
<code>--reject <LIST></code>	(0.67.0) Comma-separated filename-suffix reject list. Combines with <code>--accept</code> (URL must pass both). URLs with empty final segments pass.
<code>--continue</code>	Resume an interrupted download (wget-style auto-detect from local file size). Equivalent to <code>--continue-at -</code> .
<code>--continue-at <OFFSET></code>	Resume from BYTE offset (curl-compatible). <code>-</code> auto-detects.
<code>--spider</code>	HEAD-only check; print <code><status> <url></code> per URL and exit non-zero on any 4xx/5xx. Pairs with <code>--input-file</code> .
<code>--timestamping</code>	Skip download when local file's mtime $\geq$ server's Last-Modified. Sets If-Modified-Since.
<code>--rate <N/s N/m N/h></code>	Request rate cap. Engages with <code>--input-file</code> : at most N requests per second/minute/hour. Overridden by <code>--wait</code> .

Examples

```
# Polite batch fetch with a 2-second gap between URLs
recon --input-file urls.txt --wait 2

# Filter to image URLs, drop thumbnails
recon --input-file urls.txt --accept jpg,png --reject thumb

# Wget-style retries (5 total attempts)
recon https://api.example.com/ --tries 5

# Spider check filtered by extension
recon --input-file urls.txt --spider --accept html,htm

# Resume an interrupted download
recon https://example.com/big.iso -o big.iso --continue
```

Output control

Flag	Description
<code>-o, --output <FILE></code>	Write body to FILE (otherwise stdout).
<code>-O, --remote-name</code>	Filename = URL's last path segment (percent-decoded).
<code>--remote-name-all</code>	Apply <code>-O</code> to every URL in <code>--input-file</code> . Curl-parity for batch downloads.
<code>-J, --remote-header-name</code>	Use Content-Disposition filename when saving.
<code>--create-dirs</code>	Create parent dirs for <code>-o</code> path.
<code>-i, --include</code>	Print response headers before body.
<code>-I, --head</code>	Print headers only; no body. Implies <code>-X HEAD</code> .
<code>-s, --silent</code>	Suppress progress + verbose output.
<code>-S, --show-error</code>	Re-enable error output even when <code>-s</code> is set (curl-compat).
<code>-v, --verbose</code>	Verbose. Repeatable: <code>-v</code> , <code>-vv</code> , <code>-vvv</code> .
<code>--progress</code>	Show a progress meter when saving to a file (opt-in).

Flag	Description
<code>-#, --progress-bar</code>	<code>#</code> -character progress bar style (curl <code>-#</code> parity). Also activates the progress meter.
<code>--pretty</code>	Pretty-print JSON / XML / YAML bodies.
<code>--pretty-as <FORMAT></code>	Force pretty format (json/xml/html/yaml/csv/tsv/auto). Implies <code>--pretty</code> . Use when auto-detect picks the wrong format or there is no Content-Type.
<code>--stdin</code>	Read body from stdin instead of making an HTTP request. Runs the post-fetch pipeline (pretty, <code>--output-charset</code> , <code>-o</code>) over the piped input. Mutually exclusive with a URL.
<code>--from-clipboard</code>	Read body from system clipboard (no HTTP request). Mutex with <code>--stdin</code> and URL.
<code>--to-clipboard</code>	Write output to system clipboard. Mutex with <code>-o</code> and <code>--editor</code> . UTF-8 text only.
<code>--clipboard [<DIR>]</code>	Use clipboard for I/O. <code>DIR</code> = <code>in</code> / <code>out</code> / <code>both</code> . Bare form auto-resolves direction from context.
<code>--no-pretty</code>	Disable auto-pretty (even when content-type suggests it).
<code>-w, --write-out <FORMAT></code>	Print a curl-compatible summary after the body. See Write-out format .
<code>--editor [<CMD>]</code>	Open response body in <code>\$EDITOR</code> after the request. URL-shaped next token (contains <code>://</code>) is treated as the positional URL — <code>recon --editor https://example.com</code> works without the <code>=</code> form.
<code>--json</code> (output-side inferred)	Combined with <code>--pretty</code> , forces JSON pretty-printing regardless of content-type.

Examples

```

# Save to file
recon https://example.com/favicon.ico -O                                # →
favicon.ico
recon https://example.com/api/users.json -o users.json
recon https://x.example.com/deep/path/file.txt --create-dirs -o
downloads/x/file.txt

# Batch download – save every URL to its own file (--remote-name-all,
0.73.0)
recon --input-file urls.txt --remote-name-all
recon --input-file urls.txt --remote-name-all --output-dir ./downloads/ --
rate 2/s

# Hash-style progress bar (curl -#, 0.73.0)
recon https://example.com/big.zip -O -#                                # # bar
instead of default
recon --input-file urls.txt --remote-name-all -#                      # combined
with batch

# Inspect headers
recon -I https://example.com/                                           # HEAD
recon -i https://api.example.com/v1/status                             # headers +
body
recon -i -I https://example.com/                                       # forced
HEAD with body suppressed

# Pretty-print
recon https://api.example.com/large.json --pretty
recon https://api.example.com/feed.xml --pretty                        #
XML also supported
curl -s https://api.example.com/blob.yaml | recon --pretty            #
stdin auto-detected

# --stdin: prettify a payload without making any HTTP request
recon --stdin --pretty                                                # auto-detect
format from body
recon --stdin --pretty-as json                                         # force JSON
recon --stdin --pretty --pretty-as xml -o pretty.xml < raw.xml      #
write to file

# Auto-detected stdin: --stdin is optional when piping
echo '{"a":1}' | recon --pretty
cat data.json | recon --pretty-as json -o pretty.json

# Clipboard I/O – native clipboard read/write without shell pipes
recon --clipboard --pretty-as json                                     # read from
clipboard, auto-resolves as input
recon --clipboard both --pretty-as json                               # prettify in place

```

```
(read from, write back to)
recon https://api.example.com/data --to-clipboard      # fetch URL, copy
result to clipboard
recon --from-clipboard --editor vim                    # open clipboard body
in editor

# --prettify-as: force the format on a real HTTP response
recon https://api.example.com/data --prettify-as json  # implies --
prettify
recon https://example.com/feed --prettify-as xml       # for servers that
lie about Content-Type

# Verbose progression
recon -v https://api.example.com/                      # connection + TLS + headers
recon -vv https://api.example.com/                     # plus intermediate request info
recon -vvv https://api.example.com/                   # plus raw handshake + wire dump

# Writeout
recon https://api.example.com/ -w '\n%{http_code} in %{time_total}s (%
{size_download} bytes)\n'
```

Authentication & TLS

Flag	Description
<code>-u, --user <USER:PASS></code>	HTTP Basic credentials.
<code>-k, --insecure</code>	Skip TLS cert verification.
<code>--tlsv1.2 / --tlsv1.3</code>	Force minimum TLS version.
<code>--cacert <PATH></code>	Trust an extra PEM root (on top of system roots).
<code>--capath <DIR></code>	Directory of <code>.pem / .crt / .cer</code> root certificates; each file is added as a trust root.
<code>--ca-native</code>	Disable built-in webpki roots; use the OS native trust store only.
<code>--crlfile <PATH></code>	PEM file of X.509 CRLs. Server certs in any loaded CRL are rejected at handshake. Multi-CRL bundles supported.
<code>--pinnedpubkey <HASHES></code>	Pin the server public key (curl <code>sha256//<base64></code> form, <code>;</code> -separated for multiple). SHA-256 of the leaf cert's SubjectPublicKeyInfo; the connection is rejected unless the key matches one. Enforced even

Flag	Description
	under <code>-k</code> . Public-key file paths are not supported. Builds a custom rustls config — incompatible with <code>--client-cert</code> / <code>--client-key</code> .
<code>--curves</code> <LIST>	Restrict TLS key-exchange groups, colon-separated in preference order (X25519:P-256:P-384 ; OpenSSL aliases like <code>prime256v1</code> accepted). P-521 is unavailable under the ring backend and errors. Incompatible with <code>--client-cert</code> / <code>--client-key</code> .
<code>--interface</code> <IP NAME>	Bind outgoing socket to a specific local IP or interface name (eth0, en0 on Linux/macOS via getifaddrs; Windows accepts IP literals only).
<code>--limit-rate</code> <RATE>	Throttle download. Suffixes: K, M, G.
<code>--speed-limit</code> <BYTES>	Minimum bytes-per-second (fails if rate drops).
<code>--speed-time</code> <SECS>	Window for <code>--speed-limit</code> (default 30).

Examples

```
recon -u alice:s3cr3t https://private.example.com/
recon -k https://self-signed.example.com/
recon --tlsv1.3 https://example.com/           # refuse downgrade to 1.2
recon --cacert /etc/corp-root.pem https://internal.corp/
recon --capath /etc/pki/ca-trust/source/ https://internal.corp/
recon --ca-native https://example.com/         # OS trust store only
recon --crlfile /etc/pki/tls/crls/all.pem https://example.com/ # reject
revoked certs
recon --pinnedpubkey 'sha256//BASE64HASH' https://example.com/ # public-
key pinning
recon --pinnedpubkey 'sha256//AAA...;sha256//BBB...' https://example.com/
# primary + backup pin
recon --curves X25519:P-256 https://example.com/ # restrict key-exchange
groups
recon --interface 10.0.0.5 https://example.com/ # use a specific source
IP
```

`--pinnedpubkey` / `--curves` build a custom rustls config (request has no setter for either), routed via `use_preconfigured_tls` . That path also honours `--cacert` / `--capath` / `--ca-native` / `--crlfile` / `--tlsv1.x` / `--tls-max` , `-k` , and `--client-cert` / `--client-key` (mTLS composes with pinning/curves). The pin is SHA-256 of the

leaf certificate's SubjectPublicKeyInfo and is enforced even under `-k`. P-521 is unavailable under the ring crypto backend.

Client certificates (mTLS)

Shipped in 0.54.0. Present a client certificate during the TLS handshake for mutual TLS.

Curl-compatible note. The short form `-E` matches curl. The long form is `--client-cert` in recon (curl spells it `--cert`) — `--cert` is already taken by recon's server certificate inspection mode, see [TLS certificate inspection](#). `-E /path/to/cert.pem` works identically in both tools.

Flag	Description
<code>-E, --client-cert <PATH></code>	PEM-encoded client cert. May include the key inline.
<code>--client-key <PATH></code>	PEM key (when cert is cert-only).
<code>--cert-type <PEM DER></code>	Cert format. DER errors with a conversion recipe.
<code>--key-type <PEM DER ENG></code>	Key format. ENG refused (rustls has no engine concept).
<code>--pass <PASS></code>	Placeholder for encrypted PKCS#8 passphrase. Encrypted keys refused with an <code>openssl pkcs8</code> recipe.

Examples

```
# Combined cert+key PEM
recon -E ~/keys/bundle.pem https://mtls.example.com/

# Split cert and key
recon -E client.crt --client-key client.key https://mtls.example.com/

# Decrypt encrypted PKCS#8 externally first
openssl pkcs8 -in client.key.enc -out client.key
recon -E client.crt --client-key client.key https://mtls.example.com/
```

See also: `recon --help client-cert`.

Browser fingerprint impersonation (0.77.0, opt-in)

Shipped in 0.77.0 behind the `impersonate` Cargo feature. Routes the request through `wreq` (BoringSSL) + `wreq-util` to mimic a real browser's TLS ClientHello and HTTP/2 SETTINGS frame. Useful when a server uses JA3/JA4 fingerprinting or H2-frame analysis to distinguish bots from real browsers. (`wreq` is the renamed successor to `rquest`, which was yanked from crates.io after the upstream rename; `recon-cli` migrated in 0.80.7.)

The default `recon` binary is built without this feature so the binary stays small and skips the BoringSSL build dependency. Build with `cargo build --release --no-default-features --features impersonate`, or download the `recon-impersonate` release artifact.

`--no-default-features` is required. BoringSSL (from `wreq`) and `libssh2` (from the default-on `ssh` feature, built against OpenSSL) cannot coexist in one binary — they collide at link time. Dropping default features removes `ssh`, so the impersonate variant links a single SSL backend. The trade-off: **`recon-impersonate` has no `scp://` / `sftp://` / `ssh://` support** (those schemes return a clear "not available in this build" error, and the protocol is absent from `recon --version`). SSH probing and browser-fingerprint impersonation are orthogonal use cases, so this split is by design. See GitHub issue #1.

Flag	Description
<code>--impersonate</code> <code><PROFILE></code>	Forwards to <code>wreq_util::Emulation</code> . Examples: <code>chrome_131</code> , <code>firefox_128</code> , <code>safari_17.5</code> , <code>edge_131</code> , <code>okhttp_5</code> , <code>chrome_android_131</code> , <code>safari_ios_17.4.1</code> . Hyphens accepted as a convenience (<code>chrome-131</code> $\equiv$ <code>chrome_131</code>). See <code>recon --help impersonate</code> for the full list of supported profiles.
<code>--ja3</code> <code><STRING></code>	Deferred. Reserved in the CLI; errors at runtime. JA3 strings drop signalgs and extension order, so reconstruction is lossy. Use <code>--impersonate</code> . See OUT-OF-SCOPE.md.
<code>--ja4</code> <code><STRING></code>	Deferred. Same — JA4's cipher/extension hashes are non-invertible SHA-256 truncations.
<code>--http2-fingerprint</code> <code><STRING></code>	Override the HTTP/2 fingerprint with an Akamai-format string <code>SETTINGS WINDOW_UPDATE PRIORITY PSEUDO_HEADER_ORDER</code> (e.g. <code>1:65536,3:1000,4:6291456,6:262144 15663105 0 m,a,s,p</code>). SETTINGS ids: 1 header-table-size, 2 enable-push, 3 max-concurrent-streams, 4 initial-window-size, 5 max-frame-size, 6 max-header-list-size, 8/9 vendor. Independent of the TLS layer — combine with <code>--impersonate</code> to keep a profile's TLS fingerprint while overriding H2, or use standalone (default TLS + custom H2).

V1 incompatibility list — these flags cannot combine with `--impersonate: --ciphers, --tls13-ciphers, --tlsv1.2, --tlsv1.3, --client-cert, --client-key, --cacert, --capath`. Reason: the impersonation profile owns the TLS configuration; user-supplied overrides would defeat the fingerprint.

Examples

```
# Build with the impersonate feature (drops ssh – see note above)
cargo build --release --no-default-features --features impersonate

# Impersonate Chrome 131 against an HTTPS endpoint
recon --impersonate chrome_131 https://httpbin.org/headers

# Hyphens also work
recon --impersonate chrome-131 https://example.com/

# Verify the live fingerprint against a JA3 / JA4 echo server
recon --impersonate firefox_128 https://tls.peet.ws/api/all

# Mobile and OkHttp profiles
recon --impersonate chrome_android_131 https://example.com/
recon --impersonate safari_ios_17.4.1 https://example.com/
recon --impersonate okhttp_5 https://example.com/

# Override just the HTTP/2 fingerprint (Akamai format), standalone...
recon --http2-fingerprint
'1:65536,3:1000,4:6291456,6:262144|15663105|0|m,a,s,p' https://example.com/

# ...or layered on a profile (profile TLS kept, H2 overridden)
recon --impersonate chrome_131 --http2-fingerprint
'1:65536,4:6291456|0|0|m,a,s,p' https://example.com/
```

The default (rustls-only) build rejects these flags with a clear "rebuild with `--features impersonate`" hint pointing at the `recon-impersonate` release artifact.

See also: `recon --help impersonate`.

Proxy routing

0.50.0 shipped the full proxy suite. Schemes auto-detected from the URL: `http://`, `https://`, `socks5://`, `socks5h://`.

Flag	Description
<code>-x, --proxy <URL></code>	Proxy URL. Falls back to <code>\$HTTPS_PROXY</code> / <code>\$HTTP_PROXY</code> / <code>\$ALL_PROXY</code> .

Flag	Description
<code>-U, --proxy-user <USER:PASS></code>	Basic auth for the proxy.
<code>-p, --proxytunnel</code>	Force tunneling via CONNECT even for <code>http://</code> origins (curl-compat). HTTPS already auto-tunnels via request.
<code>--noproxy <LIST></code>	Comma-separated bypass list (matches <code>\$NO_PROXY</code> semantics; <code>*</code> means bypass all).
<code>--proxy-insecure</code>	Skip cert verification on the TLS-to-proxy connection.
<code>--proxy-cacert <PATH></code>	Extra CA for the proxy connection.
<code>--proxy-capath <DIR></code>	Directory of <code>.pem / .crt / .cer</code> CAs for proxy TLS. Mirrors <code>--capath</code> .
<code>--proxy-ca-native</code>	Disable webpki roots for proxy TLS; use OS native roots. Mirrors <code>--ca-native</code> .
<code>--proxy-pass <PASS></code>	Passphrase for <code>--proxy-key</code> (HTTPS proxy mTLS). Accepted for curl parity; deferred — request 0.12 does not expose a passphrase API for proxy mTLS. Emits a runtime warning.

Examples

```
recon -x http://proxy.corp:3128 https://api.example.com/
recon -x socks5h://127.0.0.1:9050 https://example.onion/
recon -x https://proxy.example.com:8443 -U alice:s3cr3t
https://example.com/
HTTPS_PROXY=http://proxy.corp:3128 NO_PROXY='.internal' recon
https://api.example.com/

# Proxy CA configuration (0.72.0)
recon --proxy http://corp-proxy:3128 --proxy-capath /etc/pki/proxy/
https://example.com/
recon --proxy https://proxy:8443 --proxy-ca-native https://example.com/

# Script equivalent
recon --script - <<'EOF'
http("https://api.example.com/", #{
  proxy: "socks5h://127.0.0.1:9050",
  proxy_user: "alice:s3cr3t",
  noproxy: ".internal",
});
EOF
```

Unix-domain sockets

0.51.0. Route the HTTP request over a Unix socket instead of TCP. Host header and path are preserved; only the transport changes.

Flag	Description
<code>--unix-socket <PATH></code>	Socket path.

Examples

```
# Docker API
recon --unix-socket /var/run/docker.sock http://localhost/_ping
recon --unix-socket /var/run/docker.sock --prettify
http://localhost/v1.40/version
recon --unix-socket /var/run/docker.sock
http://localhost/v1.40/containers/json

# systemd-activated service
recon --unix-socket /run/app.sock http://localhost/api/v1/status
```

HSTS persistent cache

0.52.0. A curl-compatible HSTS cache that upgrades `http://` to `https://` before sending (when the host has a non-expired cache entry) and updates itself from `Strict-Transport-Security` response headers.

Flag	Description
<code>--hsts <FILE></code>	Load + update the HSTS cache at this path.

Examples

```
# Populate from an https:// response
recon --hsts ~/.recon/hsts.txt https://www.cloudflare.com/

# Future http:// requests to the same host are auto-upgraded
recon --hsts ~/.recon/hsts.txt http://www.cloudflare.com/
# * HSTS: upgrading http:// to https:// for www.cloudflare.com

# File format matches curl's --hsts:
cat ~/.recon/hsts.txt
# # HSTS cache (recon 0.58.1)
# # host expires_unix    (leading . = includeSubDomains)
# .www.cloudflare.com 1808493843
```

IPFS / IPNS

0.49.0. `ipfs://<cid>` and `ipns://<name>` URLs are rewritten to an HTTP gateway URL before the request fires.

Flag	Description
<code>--ipfs-gateway <URL></code>	Gateway base URL. Default <code>https://ipfs.io</code> . Also read from <code>\$RECON_IPFS_GATEWAY</code> .

Examples

```
recon ipfs://QmYwAPJzv5CZsnA625s3Xf2nemtYgPpHdWEz79ojWnPbdG
recon --ipfs-gateway http://127.0.0.1:8080 ipfs://bafy... # local Kubo
node
recon ipns://example.eth # via the
default gateway
```

Cookies

Flag	Description
<code>-b, --cookiejar <PATH></code>	SQLite cookie jar at PATH (created on demand).
<code>--cookies</code>	Read-only listing of the current cookie jar.
<code>--cookie-set <NAME=VAL></code>	Set a cookie in the jar.
<code>--cookie-delete <NAME></code>	Remove by name.

Cookies persist across invocations when you pass the same `--cookiejar` path. The `.db` format is SQLite; inspect with `sqlite3` or any sqlite viewer.

Examples

```
# Login, save cookies, browse with them
recon https://example.com/login -d 'user=alice&pass=secret' --cookiejar
~/jar.db
recon https://example.com/dashboard --cookiejar ~/jar.db --prettify

# Inspect
recon --cookiejar ~/jar.db --cookies

# Inject manually
recon --cookiejar ~/jar.db --cookie-set 'session=abc123;
Domain=.example.com; Path=/'
```

DNS

Flag	Description
<code>--dns</code>	Resolve A, AAAA, MX, TXT, NS, CNAME, SOA for the URL's host.
<code>--dns-type</code> <code><LIST></code>	Comma-separated subset: A, AAAA, MX, TXT, NS, CNAME, SOA, PTR, CAA, DS, DNSKEY, HTTPS, SRV, SVCB .
<code>--dns-servers</code> <code><LIST></code>	Comma-separated resolver IPs.
<code>--dns-ipv4-addr</code> <code><IP></code>	Bind DNS queries to a specific local IPv4.
<code>--dns-ipv6-addr</code> <code><IP></code>	Same for IPv6.
<code>--dns-interface</code> <code><NAME></code>	Accepted at the CLI; not yet plumbed (see OUT-OF-SCOPE.md). Use <code>--dns-ipv4-addr</code> / <code>--dns-ipv6-addr</code> as a workaround.

Examples

```
recon example.com --dns
recon example.com --dns-type A,AAAA,MX
recon example.com --dns --dns-servers 1.1.1.1,8.8.8.8
recon example.com --dns --dns-servers 127.0.0.1:5353      # local
resolver on custom port
```


TLS certificate inspection

Standalone mode selected by `--cert` (without a `--cert <PATH>` value — it's the bool form of the flag).

Flag	Description
<code>--cert</code>	Inspect the server cert: subject, issuer, SANs, NotBefore/After, fingerprints.
<code>--cert-chain</code>	Walk the full intermediate chain.
<code>--sni</code> <code><HOSTNAME></code>	Override the SNI value sent in the handshake.

Examples

```
recon https://example.com --cert
recon https://example.com --cert --cert-chain
recon https://example.com --cert --sni alt.example.com      # request
cert for alt via SNI
```

Configuration files

recon reads two TOML configuration layers at startup and deep-merges them:

Layer	Default path
System (optional, admin- supplied)	<code>/etc/recon/config.toml</code> (Linux). On macOS, the resolver searches <code>\$HOMEBREW_PREFIX/etc/recon/config.toml</code> , <code>/opt/homebrew/etc/recon/config.toml</code> , <code>/usr/local/etc/recon/config.toml</code> , <code>/etc/recon/config.toml</code> — first existing match wins.
User (per- user overrides)	<code>~/.recon/config.toml</code>

Both layers are optional. If neither exists, recon runs with default settings.

Environment overrides

Env var	Effect
<code>\$RECON_SYSTEM_CONFIG</code>	Override the system-layer path. Accepts a file path or a directory (the directory form appends <code>config.toml</code>).

Env var	Effect
<code>\$RECON_CONFIG</code>	Override the user-layer path. Same file-or-directory rule.

CLI flags

Flag	Behavior
<code>--disable / -q</code>	Skip both config layers entirely.
<code>--no-system-config</code>	Skip only the system layer.
<code>--no-user-config</code>	Skip only the user layer.
<code>--show-config-paths</code>	Print which file each layer resolved to plus the env vars that influenced the decision, then exit.

Skip flags always win over env-var overrides — passing `--no-system-config` with `$RECON_SYSTEM_CONFIG` set silently skips the system layer.

Deep-merge rules

When both layers are present, the user layer is merged onto the system layer:

- Tables merge recursively.
- Leaves (strings, integers, booleans, datetimes) are replaced by the overlay.
- Arrays are replaced by the overlay (no concatenation — replicating systemd / sshd drop-in behavior).
- Type clashes (table vs. leaf) resolve with the overlay winning silently; downstream serde catches schema errors.

Worked example

```
# /etc/recon/config.toml (system)
[editor]
default = "vim"
[editor.aliases]
v = "vim"
nv = "nvim"

[gh.accounts]
"shared@example.com" = "shared-gh"

[ai.backends.work]
cmd = "/opt/claude"

[netstatus]
probes = ["dns://example.com", "tcp://example.com:443"]
```

```
# ~/.recon/config.toml (user)
[editor]
default = "zed"
[editor.aliases]
v = "vimnoremap"

[gh.accounts]
"me@home" = "personal-gh"

[ai.backends.scratch]
cmd = "claude"
```

```
# Effective config after merge
[editor]
default = "zed" # leaf replaced
[editor.aliases]
v = "vimnoremap" # leaf replaced
nv = "nvim" # base retained

[gh.accounts]
"shared@example.com" = "shared-gh" # base retained
"me@home" = "personal-gh" # overlay added

[ai.backends.work]
cmd = "/opt/claude"
[ai.backends.scratch]
cmd = "claude"

[netstatus]
probes = ["dns://example.com", "tcp://example.com:443"]
```

Sections

The same `config.toml` carries several feature configurations:

- `[netstatus]` — see [Ping](#), [traceroute](#), [netstatus](#)
- `[editor]`, `[editor.aliases]` — see [Meta flags](#)
- `[sampledata.*]` — see [Sample data](#)
- `[ai]`, `[ai.backends.*]` — see [ai backends](#)
- `[gh.accounts]` — used by the `gh()` script binding for auto-account-switch (see [gh wrapper](#))

Aliases / compatibility modes

Aliases are named groups of short→long flag rewrites stored in `~/.recon/config.toml` under `[aliases.<name>]` tables. Selected via `--alias <name>` on the command line, or via a session-wide default:

```
[aliases]
default = "wget"

[aliases.wget]
"-r" = "--recursive"
"-l" = { long = "--level", takes_value = true }

[aliases.mine]
"-J" = "--json"
```

Flag	Description
<code>--alias <NAME></code>	Select an alias section. Bundled names: <code>curl</code> (no-op), <code>wget</code> .

Two bundled sections ship with recon: `curl` (empty, since recon's short flags already match `curl` by policy) and `wget` (the ~11 letters where `wget` and `curl` disagree). User entries deep-merge on top of bundled entries per key — defining `[aliases.wget] "-J" = "--json"` adds `-J` without touching the rest of the `wget` map.

Aliases are namespace plumbing, not feature plumbing. `--alias wget -r` rewrites to `--recursive`; if recon doesn't yet implement `--recursive`, clap reports "unknown flag" — exactly the right failure mode.

`-q/--disable` (skip all config) suppresses alias resolution along with the rest of the config system.

Ping, traceroute, netstatus

Stand-alone modes.

Flag	Description
<code>--ping <HOST></code>	TCP ping by default, ICMP with <code>--icmp</code> .
<code>--icmp</code>	Force ICMP (needs raw sockets; root on Linux, CAP_NET_RAW elsewhere).
<code>--ping-count <N></code>	Packets (default 4).
<code>--traceroute <HOST></code>	Same as <code>--ping HOST --traceroute</code> .
<code>--max-hops <N></code>	Hop limit (default 30).
<code>--netstatus</code>	Run the probe bundle defined in <code>~/.recon/config.toml</code> (layered — see Configuration files) [<code>netstatus</code>] .

Examples

```
recon --ping 8.8.8.8
recon --ping example.com --ping-count 10
recon --ping example.com:443 --icmp           # ICMP with explicit port hint
recon --traceroute 8.8.8.8 --max-hops 20
recon --netstatus                             # connectivity sweep
```

Email protection

A single URL can be probed for its whole email-auth stack in one invocation.

Flag	Description
<code>--spf</code>	SPF record validation (recursive include/redirect, lookup limit)
<code>--dmarc</code>	DMARC policy + record
<code>--dkim <SELECTOR></code>	DKIM for the given selector. Repeatable.
<code>--mta-sts</code>	MTA-STS policy
<code>--bimi [SELECTOR]</code>	BIMI record + optional VMC / SVG fetch
<code>--tls-rpt</code>	SMTP TLS-RPT record

Examples

```
# Full sweep
recon example.com --spf --dmarc --dkim selector1 --mta-sts --bimi --tls-rpt

# Single checks
recon example.com --spf
recon example.com --dkim selector1 --dkim selector2
recon example.com --bimi default           # explicit selector
```

SMTP

Stand-alone SMTP client. Send mail or just probe capabilities.

Flag	Description
<code>--mail-from</code> <code><ADDR></code>	Envelope sender. Required for send.
<code>--mail-to</code> <code><ADDR></code>	Envelope recipient. Repeatable.
<code>--mail-subject</code> <code><STR></code>	Subject header. Default "recon SMTP test".
<code>--mail-body</code> <code><STR></code>	Body. Accepts <code>@file</code> , <code>@-</code> .
<code>--mail-header</code> <code><H: V></code>	Extra header. Repeatable.
<code>--smtp-auth</code> <code><USER:PASS></code>	AUTH PLAIN → LOGIN chain.
<code>--smtp-helo</code> <code><NAME></code>	HELO/EHLO name. Default <code>recon.local</code> .
<code>--no-starttls</code>	Don't negotiate STARTTLS.
<code>--dkim-key</code> <code><PATH></code>	DKIM-sign outbound with this PEM key.
<code>--dkim-selector</code> <code><SEL></code>	DKIM selector.
<code>--dkim-domain</code> <code><DOM></code>	DKIM domain (defaults to <code>--mail-from</code> domain).
<code>--mail-auth</code> <code><ADDR></code>	Append <code>AUTH=<ADDR></code> to MAIL FROM (RFC 4954). Accepted but currently emits a warning — lettre 0.11 limitation, see OUT-OF-SCOPE.md.

Examples

```
# Probe only
recon smtp://mail.example.com:25

# Send
recon smtp://smtp.example.com:587 --mail-from me@example.com --mail-to
you@example.com \
    --mail-subject 'Hi' --mail-body 'Test' --smtp-auth me:s3cr3t

# With DKIM signing
recon smtp://smtp.example.com:587 --mail-from me@example.com --mail-to
you@example.com \
    --dkim-key ~/.dkim/default.pem --dkim-selector default
```

Mail retrieval (POP3, IMAP)

Flag	Description
<code>--stls</code>	POP3: upgrade via STLS after CAPA.
<code>--imap-peek</code>	IMAP: use BODY.PEEK (don't flip \Seen).

URLs: `pop3://`, `pop3s://`, `imap://`, `imaps://`.

Examples

```
recon pop3s://alice:s3cr3t@pop.example.com/
recon pop3://alice:s3cr3t@pop.example.com/ --stls
recon imaps://alice:s3cr3t@imap.example.com/INBOX
recon imaps://alice:s3cr3t@imap.example.com/INBOX/3      # fetch message
3
recon imaps://alice:s3cr3t@imap.example.com/INBOX --imap-peek
```

File transfer

0.47.0 shipped FTP/FTPS, SFTP, TFTP, Gopher. 0.71.0 plumbed all previously-stubbed flags through to their underlying protocol modules.

Protocol	URL scheme	Notes
FTP / FTPS	<code>ftp://</code> , <code>ftps://</code>	Anonymous by default; userinfo in URL or <code>-u</code> .
SFTP	<code>sftp://</code>	SSH-backed. <code>--privkey</code> for a specific key.

Protocol	URL scheme	Notes
TFTP	tftp://	RFC 1350. <code>--tftp-blksize</code> for RFC 2348 block-size negotiation.
Gopher	<code>gopher://</code> , <code>gophers://</code>	Selector fetch.

FTP flags

Flag	Description
<code>--list-only</code>	Use <code>NLST</code> instead of <code>LIST</code> (filenames only).
<code>-Q, --quote <CMD></code>	Run an FTP command before listing (repeatable).
<code>--ftp-skip-pasv-ip</code>	Use control-channel IP for PASV data, ignoring server-advertised PASV IP.
<code>--ftp-pasv / --disable-epsv / --disable-eprt</code>	Confirm passive mode (suppaftp 6 default).

TFTP flags

Flag	Description
<code>--tftp-no-options</code>	Confirm vanilla RFC 1350 mode (no RFC 2347 options).

Examples


```

# FTP – listing vs fetch
recon ftp://ftp.example.com/                                     #
directory listing
recon ftp://ftp.example.com/incoming/file.bin -o local.bin      #
retrieve
recon ftp://alice:s3cr3t@ftp.example.com/                       #
authenticated
recon ftp://ftp.example.com/ --list-only                        # NLST
(filename only)
recon ftp://ftp.example.com/ -Q 'SITE CHMOD 755 pub'           # pre-
transfer command
recon ftp://ftp.example.com/ --ftp-skip-pasv-ip                 # NAT-
friendly PASV

# SFTP
recon sftp://me@example.com/home/me/data.csv -o data.csv
recon sftp://me@example.com:2222/srv/ --privkey ~/.ssh/ci_rsa

# TFTP
recon tftp://tftp.example.com/pxelinux.cfg/default
recon tftp://tftp.example.com/boot.img -o boot.img --tftp-blksize 8192
recon tftp://tftp.example.com/config --tftp-no-options          # RFC
1350 vanilla mode

# Gopher
recon gopher://gopher.floodgap.com/
recon gopher://gopher.floodgap.com/0/gopher/proxy

```

Other protocols

SSH, SCP

Flag	Description
<code>--pubkey <PATH></code>	Path to SSH public key file (alias for <code>--ssh-pubkey</code>).

```

recon ssh://me@example.com -- uname -a                          # SSH
exec; prints stdout
recon scp://me@example.com/etc/motd -o motd                     # SCP
fetch
recon ssh://me@example.com --pubkey ~/.ssh/id_ed25519.pub -- whoami #
explicit public key

```

Telnet

```
recon telnet://router.example.com:23
```

LDAP

```
recon ldap://ldap.example.com -H 'Bind: cn=admin,dc=example,dc=com'
recon ldaps://ldap.example.com # implicit-TLS
```

WebSocket

```
recon wss://echo.websocket.events # handshake + Ping/Pong
```

RTSP

```
recon rtsp://camera.example.com/stream1 # OPTIONS
```

Dict (RFC 2229)

```
recon dict://dict.org/d:recon
```

NTP, Memcached, Redis, MQTT

```
recon ntp://pool.ntp.org # clock offset + rtt
recon memcache://cache.example.com/?stats # text-protocol stats
recon redis://cache.example.com/ # PING
recon redis://cache.example.com/ -X INFO # arbitrary RESP command
recon mqtt://broker.example.com --mqtt-topic status --mqtt-message 'hello'
recon mqtt://broker.example.com --mqtt-user alice --mqtt-pass s3cr3t ...
```

Source comparison

0.53.0 shipped `--compare`. Diff two sources (URL / path / `-` stdin). HTTP(S) sources flow through the full request pipeline so every request flag (`-H`, `-u`, `-L`, `-k`, cookies, proxy, HSTS) applies.

Flag	Description
<code>--compare <A> </code>	Two sources.
<code>--compare-format <FMT></code>	<code>unified</code> (default), <code>summary</code> , <code>sxs</code> .
<code>--compare-context <N></code>	Unified context lines (default 3).

Exit codes follow GNU `diff`: 0 identical, 1 differ, 2+ source-load error.

Examples

```
recon --compare one.json two.json
recon --compare https://api.example.com/v1/status ./baseline.json
curl -s https://a/ | recon --compare - ./baseline.json
recon --compare a.txt b.txt --compare-format summary
recon --compare a.txt b.txt --compare-format sxs
recon --compare a.json b.json --compare-context 5
```

Document conversions

0.58.0 shipped markdown / HTML / PDF. See also `--help docs`.

Since 0.101.0, the default md→PDF engine is **embedded typst** — a Chrome-free, pure-Rust renderer. It gives A4 page size, a numbered TOC (`#outline`), footer page numbers (front matter unnumbered, body arabic), native UTF-8 metadata, and full GFM, with no external dependency. It does **not** support CSS or raw HTML: `--doc-css`, `--no-default-css`, `--unsafe-html`, and raw HTML in the source all error under typst with a message pointing to `--pdf-engine chrome`. Use `--pdf-engine chrome` to opt back into the legacy agent-browser path for CSS-styled output, `--unsafe-html passthrough`, or an HTML cover page. `--doc-subject` warns under typst (typst has no PDF Subject field). `--html-to-pdf` is always chrome. The

marker works on both engines; `<!-- toc -->` places the table of contents.

Flag	Description
<code>--md-to-html <SRC></code>	Markdown → HTML via comrak. Output <code>-o PATH</code> or stdout.
<code>--md-to-pdf <SRC></code>	Markdown → PDF. <code>-o PATH</code> required. Default engine: embedded typst (Chrome-free).
<code>--html-to-pdf <SRC></code>	HTML → PDF (via agent-browser; always Chrome).
<code>--pdf-engine</code> <code><typst chrome></code>	md→PDF engine. <code>typst</code> (default, Chrome-free) or <code>chrome</code> (legacy agent-browser; required for CSS / <code>--unsafe-html</code> / cover-HTML). md→pdf only; <code>--html-to-pdf</code> is always <code>chrome</code> .
<code>--page-size <SIZE></code>	Page size for typst: <code>a4</code> (default), <code>a3</code> , <code>a5</code> , <code>letter</code> , <code>legal</code> , or <code>WxH</code> (e.g. <code>420mmx594mm</code>). Errors on <code>--pdf-engine chrome</code> .
<code>--cover</code>	Generate a title page from <code>--doc-*</code> metadata (typst).
<code>--cover-template</code> <code><FILE></code>	Custom typst cover snippet; metadata injected as <code>#let title/subtitle/author/version/date</code> . Implies a cover (typst).
<code>--doc-subtitle <STR></code>	Subtitle on the generated cover page (typst).
<code>--doc-version <STR></code>	Version string on the generated cover page (typst).
<code>--doc-date <STR></code>	Date string on the generated cover page (typst).
<code>--no-page-numbers</code>	Disable footer page numbers (typst; on by default — front matter unnumbered, body arabic).
<code>--html-to-text <SRC></code>	HTML → readable wrapped text (file / URL / stdin); lynx/w3m-style view.
<code>--render</code>	Render <code>text/html</code> HTTP responses as text inline; non-HTML passes through.
<code>--render-color</code> <code><auto always never></code>	ANSI styling of rendered text (auto-on when stdout is a TTY).
<code>--render-no-links</code>	Drop link surfacing in rendered text: in plain output, the <code>[N]</code> markers and the trailing URL reference list (anchor text stays inline); in coloured output, the inline link styling. Default keeps links.
<code>--width <N></code>	Wrap column for rendered output (default: terminal width on a TTY, else 80).
<code>--toc</code>	Inject a linkable TOC.
<code>--toc-depth <N></code>	Include headings up to <code>H N</code> (default 3).

Flag	Description
<code>--toc-title <STR></code>	TOC heading (default "Contents").
<code>--doc-title <STR></code>	<code><title></code> + PDF metadata title.
<code>--doc-author <STR></code>	Author field in PDF document properties.
<code>--doc-subject <STR></code>	Subject field in PDF document properties.
<code>--doc-keywords <STR></code>	Keywords field in PDF document properties (comma-separated).
<code>--doc-css <PATH></code>	Inline a custom stylesheet. Requires <code>--pdf-engine chrome</code> (errors on typst).
<code>--no-default-css</code>	Skip bundled default CSS. Requires <code>--pdf-engine chrome</code> (errors on typst).
<code>--gfm</code>	Enable GitHub-flavored extensions (tables, task lists, strikethrough, autolinks, footnotes, tagfilter).
<code>--unsafe-html</code>	Allow raw HTML passthrough (comrak's <code>unsafe_</code> flag). Needed for cover pages and explicit <code><div class="page-break"></code> markers. Off by default; assume the markdown is trusted when on.
<code>--page-break-on-h1</code>	Start a new PDF page before every top-level <code>#</code> heading except the first. No visible effect in HTML output (Chrome's printToPDF honours the <code>break-before: page</code> rule).

Examples

```

recon --md-to-html README.md --toc --gfm -o README.html

# typst is the default md→PDF engine: A4, numbered TOC, footer page numbers
recon --md-to-pdf book.md -o book.pdf
recon --md-to-pdf book.md --page-size letter -o book.pdf

# Auto cover page from metadata (typst)
recon --md-to-pdf book.md --cover \
    --doc-title 'My Book' --doc-subtitle 'A Manual' \
    --doc-author 'Alice' --doc-version 1.2.0 --doc-date 2026-06-18 -o
book.pdf

# Custom typst cover template (receives #let
title/subtitle/author/version/date)
recon --md-to-pdf book.md --cover --cover-template cover.typ --doc-title
Book -o book.pdf

recon --md-to-pdf CHANGELOG.md --toc --gfm --doc-title 'recon release
notes' -o changelog.pdf
recon --html-to-pdf report.html -o report.pdf
curl -s https://example.com/doc.md | recon --md-to-html - --toc -o doc.html

# Legacy Chrome engine – required for CSS / --unsafe-html / cover-HTML
recon --md-to-pdf notes.md --pdf-engine chrome --doc-css print.css -o
notes.pdf
recon --md-to-pdf notes.md --pdf-engine chrome --no-default-css --doc-css
corp.css -o notes.pdf

# Book-style cover page + chapter breaks (Chrome engine)
recon --md-to-pdf book.md --pdf-engine chrome --toc --gfm --unsafe-html --
page-break-on-h1 \
    --doc-title 'My Book' -o book.pdf

# Full PDF metadata – verifiable via pdftinfo
recon --md-to-pdf report.md \
    --doc-title 'Q1 Results' \
    --doc-author 'Alice Smith' \
    --doc-subject 'Quarterly financial report' \
    --doc-keywords 'finance, Q1, 2026' \
    -o report.pdf
# pdftinfo report.pdf | grep -E '(Title|Author|Subject|Keywords)'

```

Non-ASCII metadata is fully supported: `--doc-author 'Björk'`, `--doc-subject 'Café résumé'`, etc. are written as UTF-16BE PDF strings (BOM-prefixed) so conformant readers decode them correctly — ASCII values use the plain literal form.

Rendering HTML as text

```
# Read a page as text (lynx/w3m style)
recon --render https://example.com

# Convert a local file, wrap to 60 columns
recon --html-to-text page.html --width 60 -o page.txt

# Drop the URL footnote list – keep just the anchor text
recon --render --render-no-links https://example.com
```

`--render` only transforms `text/html` responses; other content types pass through unchanged. `<a href>` links become `[N]` footnotes with a reference list (plain mode); in colour mode on a TTY, links are styled inline. ANSI styling auto-enables on a TTY — control it with `--render-color`. Pass `--render-no-links` to drop the footnote list (plain) or the inline link styling (coloured) — the anchor text stays either way. The same toggle is available as the `no_links` opts key on the `html_to_text()` script binding.

Cover pages

With `--unsafe-html`, raw HTML passes through the markdown parser verbatim. The bundled CSS styles `<div class="cover">` as a full-page centered block with an automatic page break after:

```
<div class="cover">

# My Document

<div class="subtitle">An illustrated guide</div>

<hr>

<div class="version">Version 1.0</div>
<div class="date">2026-04-24</div>
<div class="author">Alice Example</div>

</div>

# First chapter

Body text...
```

The cover block uses a `min-height: 90vh` flex container with centered content. `.subtitle`, `.version`, `.date`, `.author`, and `.meta` child classes pick up smaller, muted styling. A horizontal rule becomes a narrow divider.

Page breaks

Two routes:

1. `--page-break-on-h1` — every top-level `#` heading starts a new PDF page (except the first, which opens the document). No markdown changes required.
2. **Explicit markers (needs `--unsafe-html`)** — embed one of:

```
<div class="page-break"></div>
<hr class="page-break">
```

anywhere in the markdown. The CSS `break-after: page` kicks in.

The CSS cover block and `page-break-*` CSS described above apply to the **Chrome engine** (`--pdf-engine chrome`). Under the default typst engine, use `--cover` / `--cover-template` for the title page and the

marker for breaks. With `--pdf-engine chrome`, Chrome's printToPDF renders both `--md-to-pdf` and `--html-to-pdf`, so any CSS that speaks `break-before`, `break-after`, or `page-break-*` works; `@page` rules for size and margins (in the bundled default CSS) are honored too — override via `--doc-css` or inline `<style>` with `--unsafe-html`.

PDF page export

Flag	Value	Description
<code>--export-pdf-page</code>	<code><PAGE> <PDF></code>	Render the 1-indexed PAGE of PDF as a raster image. Requires <code>pdftoppm</code> (poppler-utils).
<code>--pdf-viewport</code>	<code><WxH></code>	Target image box in pixels. Default <code>1024x1366</code> . Aspect is preserved — this is an upper-bound box.
<code>--pdf-scale</code>	<code><N></code>	Density multiplier (≥ 1). Default <code>2</code> . Final image fits within $W*N \times H*N$ px.
<code>--pdf-quality</code>	<code><0-100></code>	JPEG/WEBP quality. Default <code>90</code> .
<code>--pdf-format</code>	<code><png jpeg webp></code>	Override output format inference.

Examples:

```
recon --export-pdf-page 1 docs/MANUAL.pdf
# → page-1.png in CWD

recon --export-pdf-page 3 report.pdf -o cover.jpg --pdf-quality 75
recon --export-pdf-page 1 doc.pdf -o cover.webp --pdf-viewport 1920x2715
```

Barcode encoding & decoding

Flag	Description
<code>--encode <FORMAT></code>	qr, datamatrix, code128, code39, ean13, upca, aztec, pdf417.
<code>--encode-format <FMT></code>	ascii (default), svg, png. Inferred from <code>-o</code> extension.
<code>--from-file <PATH></code>	Encode input from file (mutually exclusive with positional).
<code>--encode-list</code>	List all supported formats.
<code>--qr-level <L M Q H></code>	QR error-correction level (default M).

Flag	Description
<code>--encode-hints</code> <code><KEY=VAL></code>	rxing encoder hint (repeatable). Applies to aztec / pdf417. Keys: <code>charset</code> , <code>eclevel</code> , <code>aztec-layers</code> , <code>pdf417-compact</code> , <code>pdf417-compaction</code> , <code>pdf417-auto-eci</code> , <code>margin</code> .
<code>--hrt</code> / <code>--no-hrt</code>	Human-readable text under 1D barcodes. Default on for EAN/UPC, off for Code128/39. Renders to ASCII, SVG, and PNG (PNG uses the bundled DejaVu Sans Mono font).
<code>--decode</code> <code><IMAGE></code>	Scan a PNG/JPEG/WebP for any supported format. <code>-</code> for stdin.
<code>--decode-hints</code> <code><LIST></code>	Comma-separated format restriction.
<code>--decode-all</code> <code><IMAGE></code>	Scan for every barcode in the image; one line per detection.

Examples

```
# Encode
recon --encode qr -o qr.png 'https://example.com'
recon --encode qr --qr-level H -o durable.png 'sticker text'
recon --encode pdf417 -o id.png 'stacked linear code'
recon --encode aztec -o ticket.png 'transit ticket'
recon --encode ean13 --encode-format svg '4006381333931' > product.svg
recon --encode qr --from-file msg.txt -o qr-of-file.png

# Decode
recon --decode ticket.png                # → aztec<TAB>transit
ticket
cat code.png | recon --decode -
recon --decode mystery.png --decode-hints qr,datamatrix
recon --decode bottle.jpg --decode-hints ean13
recon --decode-all sheet.png            # every code in the
image
```

`--encode-hints` — Aztec / PDF417 tuning

`--encode-hints KEY=VAL` exposes the rxing `encode_with_hints` API for the two formats recon routes through rxing (Aztec and PDF417). The flag is repeatable; unknown keys error, as do hints applied to non-rxing formats (qr, datamatrix, code128, code39, ean13, upca).

Supported keys:

Key	Applies to	Value
charset	aztec, pdf417	Character set name driving the ECI selection (e.g. UTF-8 , Shift_JIS , ISO-8859-1).
eclevel	aztec, pdf417	Aztec: minimum % of EC words. PDF417: integer 0..8 (higher = more redundancy).
aztec-layers	aztec	-4..-1 for compact mode, 0 for auto, 1..32 for full mode.
pdf417-compact	pdf417	true / false . Use PDF417 compact mode.
pdf417-compaction	pdf417	Compaction mode name (e.g. TEXT , BYTE , NUMERIC).
pdf417-auto-eci	pdf417	true / false . Auto-insert ECIs for non-Latin-1 input.
margin	aztec, pdf417	Quiet-zone margin in pixels.

```
# Compact Aztec (negative layer count)
recon --encode aztec --encode-hints aztec-layers=-2 'compact transit
ticket'

# PDF417 with maximum EC redundancy and explicit UTF-8 ECI
recon --encode pdf417 \
    --encode-hints eclevel=8 \
    --encode-hints charset=UTF-8 \
    -o id.svg 'license payload'

# Aztec with Shift_JIS ECI for a Japanese payload
recon --encode aztec --encode-hints charset=Shift_JIS '  '
```

HRT (human-readable text under 1D barcodes)

EAN-13 and UPC-A get HRT by default; Code128 and Code39 don't unless you pass `--hrt` . Implemented for ASCII, SVG, and PNG output. PNG rasterization uses a bundled DejaVu Sans Mono font (Bitstream Vera + DejaVu license; see `assets/fonts/LICENSE-DejaVu.txt`). The PNG HRT band is floored at 22 pixels tall so the text stays legible even at 1D's 2-pixel-per-module bar width.

```
recon --encode ean13 --encode-format svg '4006381333931' -o retail.svg
recon --encode code128 --hrt --encode-format svg 'SHIP-4711' -o box.svg
recon --encode ean13 --no-hrt '4006381333931' -o bare.svg
recon --encode ean13 '4006381333931' -o retail.png          # PNG HRT
(0.93.0+)
recon --encode code128 --hrt 'SHIP-4711' -o box.png
```

Hashing

Flag	Description
<code>--hash</code> <code><ALGO></code>	md5, sha1, sha224, sha256, sha384, sha512, sha3_256, sha3_512, blake3.
<code>--hash-list</code>	List supported algorithms.

Examples

```
recon --hash sha256 /etc/hosts
recon --hash sha256 https://example.com/file.iso          # hash a URL
cat payload.bin | recon --hash blake3 -
echo -n "hello" | recon --hash md5 -
```

Compression & archiving

Flag	Description
<code>--compress <ALGO></code>	gzip, deflate, brotli, zstd, bzip2, lz4, xz, snap.
<code>--decompress <ALGO></code>	Same set; auto-detect via magic bytes if ALGO is <code>auto</code> .
<code>--compress-list</code>	List supported algorithms.
<code>--compression-level</code> <code><N></code>	1..22 (depends on algo).
<code>--archive <DEST></code>	Create zip/tar/tar.gz/tar.xz/tar.bz2. Positional args after DEST are sources.
<code>--extract <SRC></code>	Extract an archive. <code>-o DIR</code> to choose destination.

Examples

```

recon --compress gzip /etc/hosts -o hosts.gz
recon --decompress auto hosts.gz -o hosts.plain
recon --compress zstd --compression-level 19 big.bin -o big.zst

recon --archive backup.tar.gz ~/Documents ~/src
recon --extract backup.tar.gz -o restore/
recon --extract release.zip

```

Encryption

age-based + PGP shell-out.

Flag	Description
<code>--encrypt</code>	Encrypt stdin / file to age format. Requires <code>--recipient</code> or <code>--pgp-recipient</code> .
<code>--decrypt</code>	Decrypt age / PGP. Needs <code>--identity</code> or a configured GPG keyring.
<code>--encrypt-keygen</code>	Generate a new age keypair.
<code>--recipient <KEY></code>	Age recipient (X25519 or ssh-ed25519). Repeatable.
<code>--identity <PATH></code>	Age identity file for decryption.
<code>--passphrase</code>	Use passphrase-based (scrypt) encryption. Mutually exclusive with <code>--recipient</code> .
<code>--armor</code>	Armored (base64 PEM) output.
<code>--pgp-recipient <EMAIL FPR></code>	PGP recipient (shells out to <code>gpg</code>).

Examples

```

# Generate
recon --encrypt-keygen > age.key          # public key printed to stderr

# Encrypt
recon --encrypt --recipient age1abc... README.md -o README.md.age
recon --encrypt --passphrase secret.txt -o secret.txt.age --armor

# Decrypt
recon --decrypt --identity age.key README.md.age -o README.md

# PGP
recon --encrypt --pgp-recipient alice@example.com msg.txt -o msg.txt.pgp

```

Check digits

Flag	Description
<code>--checkdigit <ALGO></code>	Verify an input.
<code>--checkdigit-create <ALGO></code>	Compute and append the check digit.
<code>--checkdigit-list</code>	List supported algorithms (70+).

Notable algorithms shipped in 0.61.0:

- **Latin-American + Australian + Mexican tax IDs:** `br_cpf`, `br_cnpj`, `ar_cuit`, `ar_cuil`, `cl_rut`, `pe_ruc`, `au_abn`, `mx_rfc`. All except ABN support `--checkdigit-create`; ABN's mod-89 algorithm has no single-digit inverse.
- **110+ year warnings** on Nordic + Bulgarian personal IDs: `cpr` (Denmark), `henkilotunnus` (Finland), `fodselsnummer` (Norway), `bg-egn` (Bulgaria). Same idiom as the Swedish `personnummer`: when the parsed birth year implies age ≥ 110 , the verdict's `comment` flags a likely data-entry error.

Examples

```
recon --checkdigit isbn13 '9780131103627'      # → valid
recon --checkdigit-create isbn13 '978013110362'  # → 9780131103627
recon --checkdigit-create ean13 '400638133393'  # → 4006381333931
recon --checkdigit personnummer '19900101-1239' # Swedish personal ID
recon --checkdigit-list | head -20
```

Sample data

Flag	Description
<code>--sample <NAME></code>	Emit a sample payload (faker-style).
<code>--sample-list</code>	List all named samples.

Examples

```
recon --sample-list
recon --sample json-user
recon --sample json-user -o user.json
recon --sample css                # minimal CSS
boilerplate
```

JWT tokens

Flag	Description
<code>--jwt-view <TOKEN></code>	Decode + pretty-print. No signature check.
<code>--jwt-sign <PAYLOAD></code>	Sign a JWT. Pair with <code>--jwt-secret</code> or <code>--jwt-key</code> .
<code>--jwt-validate <TOKEN></code>	Validate signature + standard claims.
<code>--jwt-secret <STR></code>	HMAC secret (HS256/384/512).
<code>--jwt-key <PATH></code>	RSA/ECDSA/Ed25519 PEM key.
<code>--jwt-alg <ALG></code>	HS256, HS384, HS512, RS256, RS384, RS512, ES256, ES384, EdDSA.

Examples

```
recon --jwt-view "$(cat token.txt)"
recon --jwt-sign '{"sub":"alice","exp":9999999999}' --jwt-secret sharedkey
--jwt-alg HS256
recon --jwt-validate "$TOKEN" --jwt-secret sharedkey
recon --jwt-sign '{"sub":"alice"}' --jwt-key private.pem --jwt-alg RS256
```

Text encoding

Flag	Description
<code>--iconv <SRC:DST></code>	Standalone <code>iconv</code> replacement.
<code>--list-charsets</code>	List supported charset labels.

Examples

```
recon --iconv utf-8:iso-8859-1 greek.txt -o greek-latin1.txt
echo 'Hello' | recon --iconv utf-8:utf-16le - -o hello.utf16
recon --list-charsets | head -20
```

Serve mode

Flag	Description
<code>--serve [ADDR]</code>	Start an HTTP server serving the cwd (default <code>127.0.0.1:8000</code>).
<code>--serve-tls [ADDR]</code>	HTTPS. Requires <code>--serve-cert</code> + <code>--serve-key</code> .
<code>--serve-cert <PATH></code>	Server cert (PEM).
<code>--serve-key <PATH></code>	Server key (PEM).
<code>--serve-sni</code> <code><HOST:CERT:KEY></code>	Per-hostname SNI mapping. Repeatable.

Examples

```
recon --serve
recon --serve 0.0.0.0:3000
recon --serve-tls 0.0.0.0:8443 --serve-cert cert.pem --serve-key key.pem
recon --serve-tls :443 --serve-sni a.example.com:a.pem:a.key --serve-sni
b.example.com:b.pem:b.key
```

Write-out format

The `-w '<FORMAT>'` flag prints a curl-compatible summary after the response. Variable names match curl's.

Variable	Description
<code>%{http_code}</code>	HTTP status code
<code>%{size_download}</code>	Body size in bytes
<code>%{size_header}</code>	Response header bytes
<code>%{size_request}</code>	Request bytes sent
<code>%{size_upload}</code>	Uploaded bytes
<code>%{speed_download}</code>	Bytes-per-second
<code>%{time_total}</code>	Total operation time (seconds, fractional)
<code>%{time_starttransfer}</code>	TTFB
<code>%{time_redirect}</code>	Time spent on redirects
<code>%{url}</code>	Final URL (after redirects)
<code>%{url_effective}</code>	Same as <code>%{url}</code>
<code>%{content_type}</code>	Response Content-Type

Variable	Description
<code>%{num_redirects}</code>	Redirect count
<code>%{remote_ip}</code>	Final peer IP
<code>%{remote_port}</code>	Peer port
<code>%{local_ip} / %{local_port}</code>	Local side
<code>%{scheme}</code>	http or https

Not yet accurate: `time_namelookup`, `time_connect`, `time_appconnect`, `time_pretransfer` render as `0.000000` — see OUT-OF-SCOPE.md.

Examples

```
recon https://example.com/ -w '\n%{http_code} %{time_total}s %
size_download}B\n'
recon -I https://example.com/ -w '%{scheme}://%{remote_ip}:%{remote_port} %
{http_code}\n'
recon -L https://short.url/x -w '%{num_redirects} hops → %
{url_effective}\n'
```

Browser automation

Flag	Description
<code>--browser-screenshot <URL></code>	Open in agent-browser, save a PNG to <code>-o PATH</code> .

Requires `agent-browser` on `$PATH`. See [Script engine → agent-browser](#) for deeper automation.

Examples

```
recon --browser-screenshot https://example.com -o example.png
```

Meta flags

Flag	Description
<code>--flags</code>	Alphabetical curl-style flag listing (<code>--help all</code> equivalent).
<code>--examples</code>	Curated examples, paged.
<code>--init</code>	Bootstrap <code>~/ .recon/</code> (script/, jars/, sni/, config.toml).

Flag	Description
<code>--editor</code>	Open response in <code>\$EDITOR</code> .
<code>--editor-cleanup</code>	Remove leftover <code>~/.recon/editor-*</code> tempfiles.
<code>--no-pager</code>	Disable paging of <code>--help</code> / <code>--examples</code> / <code>--flags</code> .

`--flags` — the quick lookup

`recon --flags` prints every flag alphabetically sorted by long name, curl's `--help all` layout:

```
(short or 4-space pad) --long <VALUE>  short description
```

Short description is capped at ~52 characters (first sentence of each flag's internal help text). Use this as the quick index when you know roughly what you want but not the flag name; follow up with `recon --help <topic>` for the long-form deep-dive.

Example:

```
$ recon --flags | head
  --age                Force the age backend, even if...
  --archive <DEST>     Create an archive
  --armor              Produce ASCII-armored output...
  --bimi <SELECTOR>    Validate the BIMI record
  --browser-screenshot <URL> Open URL in a browser and save...
  --cacert <PATH>      Path to a PEM-encoded CA certificate...
  --cert              Fetch and display the server's TLS cert
  --cert-type <PEM|DER> Format of --client-cert
-E, --client-cert <PATH> Client certificate for mTLS
  --client-key <PATH>  Private key for the client certificate
```

The listing is auto-paged through `$PAGER`. Pipe to `grep` for search:

```
recon --flags | grep -i cookie
recon --flags | grep -E '^s*-[a-zA-Z],' # only flags with short keys
```

Interactive REPL

The `--repl` flag opens an interactive Rhai prompt backed by the script engine. Every binding available in `--script` mode is available at the prompt: `http()`, `hash_sha256()`, `encrypt_*`, `sqlite_*`, and so on.

Launching

Flag	Description
<code>--repl</code>	Open the REPL. Mutually exclusive with <code>--script</code> (REPL wins if both are given).
<code>--repl-history PATH</code>	Override the history file (default <code>~/.recon/repl_history</code>).

State persists across lines: `let` bindings and `fn` definitions remain in scope until you `:reset` or exit.

Meta-commands

All meta-commands start with `:` so they never collide with valid Rhai code.

Command	Behaviour
<code>:help</code>	Print the REPL cheat sheet.
<code>:help <topic></code>	Print <code>recon --help <topic></code> content without leaving the REPL.
<code>:load <path></code>	Eval <code><path></code> in the current scope. <code>let</code> / <code>fn</code> persist. Path resolves: literal then <code>~/.recon/script/<path>[.rhai]</code> .
<code>:run <path></code>	Eval <code><path></code> in a fresh, throwaway scope. Prints the return value. REPL state untouched.
<code>:paste</code>	Enter paste mode. Lines accumulate until <code>:end</code> alone on a line. Then compile + eval once.
<code>:set <key> <val></code>	Mutate flags. Keys: <code>method</code> , <code>header</code> (append), <code>timeout</code> , <code>user-agent</code> , <code>autoprint</code> (on/off).
<code>:vars</code>	List bound variables (consts and <code>let</code> bindings) with type tag and value preview.
<code>:fns</code>	List user-defined functions from the accumulated AST chain.
<code>:reset</code>	Clear user bindings and user functions. Keeps engine, history, and the const bindings (<code>args</code> , <code>flags</code> , ...).
<code>:save <path></code>	Write this session's successful input lines to <code><path></code> with a timestamp header.
<code>:save-tidy <path></code>	Like <code>:save</code> , but appends missing <code>;</code> and drops entries that fail to compile so the saved file runs as a script.
<code>:functions [all] / :function-list</code>	List every callable registered with the engine (probes, helpers, builders). Pass <code>all</code> to also include the Rhai standard library.
<code>:history [N]</code>	Print the last <code>N</code> (default 20) inputs with 1-based indices.
<code>!:N</code>	Re-run history entry <code>N</code> .

Command	Behaviour
<code>:edit</code>	Open <code>\$EDITOR</code> (fallback <code>vi</code>) with a temp <code>.rhai</code> file. Eval the contents on save+quit.
<code>:time <expr></code>	Evaluate <code><expr></code> and print the elapsed wall-clock time.
<code>:quit / :exit</code>	Save the history file and exit with code 0.

Multi-line input

The REPL detects unbalanced braces, parentheses, and strings and shows a `...` continuation prompt until the buffer parses. Use `:paste` to force a multi-line capture when the auto-detector mis-classifies pasted content; lines accumulate until `:end` on its own line.

Autoprint

Bare expressions print their result automatically (Python/Node convention). Toggle off with `:set autoprint off` if you want script-mode silence. `let` and `fn` declarations never print.

Threading caveat

`thread_spawn` is not available in REPL mode — it requires a static AST handle that the per-line REPL doesn't have. Calls return an error. Use `--script` for threaded workflows.

Example session

```
$ recon --repl
recon REPL - :help for commands, :quit to exit
>>> let response = http("https://api.example.com/users/1")
>>> response.status
200
>>> response.body.from_json().name
"Alice"
>>> fn pretty(r) { r.body.from_json().name }
>>> pretty(response)
"Alice"
>>> :vars
    let response = #{status: 200, body: ...}
>>> :save investigation.rhai
saved 4 entries to investigation.rhai
>>> :quit
```

Part III — Script engine

recon ships an embedded [Rhai](#) interpreter. Every protocol probe, every cryptography primitive, every helper is exposed to scripts. Scripts run via `--script PATH`; `return N` from the top level becomes the process exit code.

Running scripts

```
recon --script ./probe.rhai           # run an absolute /
relative path
recon --script probe                  # falls back to
~/recon/script/probe.rhai
recon --script probe https://example.com alice  # positional args →
args[1], args[2]
```

Script lookup order:

1. Exact path (with or without `.rhai` extension).
2. `~/recon/script/<name>.rhai`.
3. If neither exists — error.

The shipped example scripts in `script/*.rhai` are copy-friendly starting points:

```
cp script/*.rhai ~/recon/script/
recon --script http example.com
```

Bare-word invocation

After `recon --init` + copying scripts to `~/recon/script/`, scripts become first-class by name:

```
recon --script tcp-echo 127.0.0.1:9000
recon --script doc-convert README.md output.pdf
recon --script client-cert ~/keys/bundle.pem https://mtls.example.com/
```

Script language (Rhai)

Rhai is a lightweight embedded script language. Key facts for newcomers:

- **Let bindings:** `let x = 42;` (mutable by default; use `const` for compile-time constants).
- **Types:** integers (i64), floats (f64), bools, strings, arrays, maps (object literals: `#{ key: value }`), Blob (Vec).
- **String interpolation:** ``Hello ${name}``.
- **Closures / function pointers:** `|a, b| { a + b }`.
- **For loops:** `for i in 0..5 { ... }` (integer ranges), `for item in array { ... }`, `for (k, v) in map { ... }`.
- **Functions:** `fn f(a) { a * 2 }`.
- **Imports:** `import "module-name" as m;` — resolved against the script's directory and `~/.recon/script/`.
- **Error handling:** `try { ... } catch(e) { ... }`.
- **Reserved words** that might surprise: `spawn` (use `thread_spawn`), `async`, `await`, `throw` (all reserved even when unused).

Useful idioms:

```
// Map with computed keys
let cfg = #{
  user: env("USER"),
  host: "api.example.com",
  timeout_ms: 5000,
};

// Chainable array methods
let evens = (0..20).filter(|n| n % 2 == 0).map(|n| n * 10);

// Early return
if !file_exists("/etc/passwd") { return 1; }
```

Types and literals

Rhai's type system is fixed (no user-defined types at runtime). Recon scripts use the following built-in types:

Type	Rust equivalent	Example literals
i64	64-bit signed integer	42, -7, 0xFF, 0b1010, 0o77
f64	64-bit float	3.14, -1.0e-5, 1_000.5
bool	Boolean	true, false
()	Unit / null	()
char	Unicode scalar	'a', '\n'
String	UTF-8 string	"hello", `hi \${name}`

Type	Rust equivalent	Example literals
Array	Heterogeneous dynamic array	<code>[1, "two", 3.0]</code>
Map	String-keyed object	<code>#{ x: 1, y: "ok" }</code>
Blob	Byte array (<code>Vec<u8></code>)	<code>Blob(), "hello".to_blob()</code>

```

let n    = 255;           // i64 – integer literal
let f    = 3.14;          // f64
let flag = true;          // bool
let nil  = ();            // unit – Rhai's null
let ch   = 'A';           // char
let s    = "hello";       // string
let tmpl = `n is ${n}`;   // interpolated string (backtick)
let arr  = [1, "two", false]; // heterogeneous array
let obj  = #{ host: "localhost", port: 8080 }; // map
let raw  = "hello".to_blob(); // Blob of UTF-8 bytes

```

Integer literals accept `_` separators (`1_000_000`) and radix prefixes `0x` (hex), `0b` (binary), `0o` (octal). Type annotations (`let x: i64 = 0;`) are optional — Rhai infers the type.

Operators

Arithmetic

Operator	Description	Example
<code>+</code>	Addition; string concatenation	<code>1 + 2</code> , <code>"a" + "b"</code>
<code>-</code>	Subtraction	<code>10 - 3</code>
<code>*</code>	Multiplication	<code>4 * 5</code>
<code>/</code>	Division (integer truncates)	<code>7 / 2 → 3</code> ; <code>7.0 / 2.0 → 3.5</code>
<code>%</code>	Remainder	<code>10 % 3 → 1</code>
<code>**</code>	Power	<code>2 ** 10 → 1024</code>
<code>-x</code>	Unary negation	<code>-x</code>

Comparison (all return `bool`)

`==`, `!=`, `<`, `<=`, `>`, `>=` — work on all numeric types and strings (lexicographic for strings).

Logical

Operator	Description
<code>&&</code>	Short-circuit AND
<code> </code>	Short-circuit OR
<code>!</code>	Logical NOT

Bitwise (integers only)

`&` (AND), `|` (OR), `^` (XOR), `~` (bitwise NOT), `<<` (left shift), `>>` (right shift).

String concatenation

`"hello " + "world"` — either operand may be a non-string and is automatically converted:
`"n=" + 42 → "n=42"`.

Null-coalescing

`value ?? default` evaluates to `default` when `value` is `()`, otherwise `value`. Useful for optional map lookups — accessing a missing key returns `()`:

```
let port    = opts["port"] ?? 80;
let charset = response["content-type"] ?? "application/octet-stream";
```

Note: `env()` always returns a `String` (empty string when unset), never `()`, so `??` does not apply to it. Use the two-argument form for env defaults: `env("HOME", "/tmp")`.

Ranges

`1..5` (exclusive — does not include 5), `1..=5` (inclusive — includes 5). Ranges are iterable and support `.filter()`, `.map()`, `.reduce()`.

Combined example

```
let x = 7;
let y = 2;
print(x + y);           // 9
print(x ** y);          // 49
print(x % y);           // 1
print(x > 5 && y < 10); // true
print((x | y) == 7);    // true (0b111 | 0b010 = 0b111)
let label = env("LABEL", "default"); // env() never returns () — use 2-arg form
print(`label is ${label}`);
let sum = (1..=10).reduce(|acc, n| acc + n, 0);
print(sum);             // 55
```


Control flow

if / else if / else

```
if x > 10 {  
    print("big");  
} else if x > 0 {  
    print("small");  
} else {  
    print("non-positive");  
}
```

If blocks are expressions and return the last evaluated value:

```
let label = if x > 0 { "positive" } else { "non-positive" };
```

Ternary operator

```
let sign = x > 0 ? "+" : "-";
```

while

```
let i = 0;  
while i < 5 {  
    print(i);  
    i += 1;  
}
```

loop

`loop { ... }` runs until `break`:

```
let attempts = 0;  
loop {  
    attempts += 1;  
    if attempts >= 3 { break; }  
}
```

for — range

```
for i in 0..5 { print(i); } // 0 1 2 3 4 (exclusive)  
for i in 1..=5 { print(i); } // 1 2 3 4 5 (inclusive)
```

for — array

```
let hosts = ["a.com", "b.com", "c.com"];
for host in hosts {
    print(host);
}
```

for — map

```
let cfg = #{ host: "a.com", port: 443 };
for (key, val) in cfg {
    print(`${key} = ${val}`);
}
```

break / continue

```
for i in 0..10 {
    if i == 3 { continue; } // skip 3
    if i == 7 { break; }    // stop before 7
    print(i);
}
```

return

`return` exits the current function (or the script when used at top level):

```
fn first_positive(arr) {
    for n in arr {
        if n > 0 { return n; }
    }
    () // not found – implicit unit return
}
```

Functions and closures

Named functions

```
fn add(a, b) { a + b }           // last expression is implicit return value

fn factorial(n) {
    if n <= 1 { return 1; }
    n * factorial(n - 1)         // recursion works
}

print(add(3, 4));                // 7
print(factorial(5));             // 120
```

Rhai hoists function definitions, so a function can be called before it appears in the file.

Closures

```
let double = |x| x * 2;
let greet  = |name| `Hello, ${name}!`;

print(double(5));      // 10
print(greet("world")); // Hello, world!

let nums  = [1, 2, 3, 4, 5];
let evens = nums.filter(|n| n % 2 == 0); // [2, 4]
let square = evens.map(|n| n * n);       // [4, 16]
print(square);
```

Closures capture outer variables by value at creation time.

Function pointers

```
fn triple(n) { n * 3 }

let fp = Fn("triple"); // function pointer by name
print(call(fp, 7));    // 21
```

Error handling

try / catch

In Rhai 1.x, `e` in a `catch` block is the error message as a string:

```
try {
    let r = http("https://httpbin.org/status/500");
    if r.status != 200 {
        throw `HTTP error: ${r.status}`;
    }
    print(r.body);
} catch(e) {
    eprint(`Error: ${e}`);
}
```

throw

`throw "message"` raises an error that propagates to the nearest `catch` block or terminates the script with a non-zero exit code:

```
fn parse_port(s) {
  let n = s.to_int();
  if n < 1 || n > 65535 { throw `invalid port: ${s}`; }
  n
}
```

assert

`assert(cond, msg)` is a recon helper that calls `throw msg` when `cond` is `false`. Prefer it over manual `if/throw` when validating invariants:

```
let r = http("https://httpbin.org/get");
assert(r.status == 200, `Expected 200, got ${r.status}`);
```

Standard built-in functions

The following are available in every recon script without any import. They come from Rhai's standard packages, augmented by recon's own helpers.

Type inspection and conversion

Expression	Description	Example result
<code>type_of(v)</code>	Runtime type name	"i64", "f64", "string", "array", "map", "bool", "()", "blob"
<code>v.to_int()</code>	Convert to i64	"42".to_int() → 42
<code>v.to_float()</code>	Convert to f64	3.to_float() → 3.0
<code>v.to_string()</code>	Convert to String	true.to_string() → "true"
<code>parse_int(s)</code>	Parse decimal string → i64	parse_int("255") → 255
<code>parse_float(s)</code>	Parse string → f64	parse_float("3.14") → 3.14

```
let v = "123".to_int();
print(type_of(v));           // "i64"
print(type_of(3.14));        // "f64"
print(type_of([1, 2]));      // "array"
print(type_of(()));          // "()"
print(42.to_float() + 0.5);  // 42.5
print(parse_int("0xFF"));    // 255
```

String methods

Strings in Rhai are immutable. Methods that appear to "modify" a string return a new one.

Method	Description
<code>.len()</code>	Character (code-point) count
<code>.is_empty()</code>	<code>true</code> if <code>""</code>
<code>.to_upper()</code>	Uppercase copy
<code>.to_lower()</code>	Lowercase copy
<code>.trim()</code>	Strip leading and trailing whitespace
<code>.trim_start()</code> / <code>.trim_end()</code>	Strip one side only
<code>.contains(sub)</code>	Substring check; returns <code>bool</code>
<code>.starts_with(s)</code> / <code>.ends_with(s)</code>	Prefix / suffix check
<code>.replace(old, new)</code>	Replace all occurrences
<code>.split(sep)</code>	Split on separator; returns array
<code>.chars()</code>	Array of single-character strings
<code>.sub_string(start)</code>	Slice from index to end
<code>.sub_string(start, len)</code>	Slice of given length
<code>.pad(len, ch)</code>	Right-pad to <code>len</code> with char <code>ch</code>
<code>.index_of(sub)</code>	First match index or <code>-1</code>
<code>.to_blob()</code>	Encode as UTF-8 <code>Blob</code>
<code>.to_int()</code>	Parse as decimal integer
<code>.to_float()</code>	Parse as float

```
let raw = " https://Example.COM/Path ";
let url = raw.trim().to_lower();           // "https://example.com/path"
print(url.contains("example"));           // true
print(url.starts_with("https"));          // true
print(url.replace("example.com", "cdn.example.com"));

let parts = "a,b,c,d".split(",");          // ["a", "b", "c", "d"]
print(parts.len());                       // 4
print(parts.join(" | "));                  // "a | b | c | d"

let host = "api.example.com";
print(host.sub_string(4, 7));              // "example"
print(host.index_of("."));                 // 3
```

Array methods

Method	Description
<code>.len()</code>	Number of elements
<code>.is_empty()</code>	<code>true</code> if array is empty
<code>.push(v)</code>	Append element (in-place)
<code>.pop()</code>	Remove and return last element
<code>.shift()</code>	Remove and return first element
<code>.unshift(v)</code>	Prepend element (in-place)
<code>.insert(i, v)</code>	Insert at index <code>i</code> (in-place)
<code>.remove(i)</code>	Remove element at index <code>i</code> and return it
<code>.reverse()</code>	Reverse in-place
<code>.contains(v)</code>	Membership test
<code>.index_of(v)</code>	First index of <code>v</code> , or <code>-1</code>
<code>.join(sep)</code>	Concatenate elements as strings with separator
<code>.map(fn)</code>	Apply function; returns new array
<code>.filter(fn)</code>	Keep elements where predicate is <code>true</code> ; returns new array
<code>.reduce(fn, init)</code>	Fold to single value
<code>.sort()</code>	In-place ascending sort
<code>.sort(fn)</code>	In-place sort with comparator <code>fn(a, b) -> i64</code>
<code>.dedup()</code>	Remove consecutive duplicate elements (in-place)
<code>.drain(range)</code>	Remove and return elements in index range

```

let nums = [5, 3, 8, 1, 9, 2];
nums.sort();                                // [1, 2, 3, 5, 8, 9]

let big  = nums.filter(|n| n > 4);           // [5, 8, 9]
let dbl  = big.map(|n| n * 2);               // [10, 16, 18]
let sum  = dbl.reduce(|acc, n| acc + n, 0);   // 44

nums.push(42);
nums.unshift(0);
print(nums.len());                          // 8
print(nums.contains(42));                   // true
print(nums.index_of(3));                    // 2 (after sort)

// Negative indices index from the end
print(nums[-1]);                           // last element

```

Map methods

Maps use string keys. Literal syntax is `#{ key: value }`. Keys with special characters or spaces must be quoted: `#{ "content-type": "json" }`.

Method / accessor	Description
<code>.len()</code>	Number of key-value pairs
<code>.is_empty()</code>	<code>true</code> if no entries
<code>.keys()</code>	Array of all key strings
<code>.values()</code>	Array of all values
<code>.contains_key(k)</code>	Key existence check
<code>.remove(k)</code>	Delete key and return its value
<code>m.key / m["key"]</code>	Read a value
<code>m.key = v / m["key"] = v</code>	Write a value

```
let hdr = #{
  "content-type": "application/json",
  accept: "application/json",
};
print(hdr.len());           // 2
print(hdr.keys());          // ["content-type", "accept"]
print(hdr.contains_key("accept")); // true

hdr["x-request-id"] = "abc-123";

for (k, v) in hdr {
  print(`${k}: ${v}`);
}

let ct = hdr.remove("content-type");
print(ct);           // "application/json"
print(hdr.len());    // 2
```

Math functions

Function	Description
<code>abs(n)</code>	Absolute value
<code>sqrt(n)</code>	Square root (returns <code>f64</code>)
<code>pow(base, exp)</code>	Power — equivalent to <code>base ** exp</code>
<code>min(a, b) / max(a, b)</code>	Minimum / maximum
<code>round(n)</code>	Round to nearest integer

Function	Description
<code>floor(n)</code> / <code>ceil(n)</code>	Floor / ceiling
<code>sin(n)</code> / <code>cos(n)</code> / <code>tan(n)</code>	Trig functions (radians)
<code>asin(n)</code> / <code>acos(n)</code> / <code>atan(n)</code>	Inverse trig (radians)
<code>ln(n)</code>	Natural logarithm (base e)
<code>log(n)</code>	Base-10 logarithm
<code>exp(n)</code>	e^n
PI	$\pi \approx 3.14159265358979$
E	Euler's number ≈ 2.71828182845905

```

print(abs(-7));           // 7
print(sqrt(144.0));       // 12.0
print(pow(2, 10));        // 1024
print(round(3.7));        // 4
print(floor(3.7));        // 3
print(ceil(3.1));         // 4
print(min(5, 3));         // 3
print(max(5, 3));         // 5

let angle = PI / 4.0;
print(sin(angle));        // ≈ 0.7071 (sin 45°)
print(exp(1.0));          // ≈ 2.7183 (same as E)

```

Ranges

Ranges are lazy integer sequences.

Expression	Meaning
<code>a..b</code>	Exclusive: integers <code>a</code> , <code>a+1</code> , ..., <code>b-1</code>
<code>a..=b</code>	Inclusive: integers <code>a</code> , <code>a+1</code> , ..., <code>b</code>

Ranges support `.filter()`, `.map()`, `.reduce()`, `.for_each()`. A range can also be used as an array slice index: `arr[start..end]`.


```
// Sum 1 to 100
let total = (1..=100).reduce(|acc, n| acc + n, 0);
print(total);    // 5050

// Collect even squares
let evens = (1..=10)
    .filter(|n| n % 2 == 0)
    .map(|n| n * n);
print(evens);    // [4, 16, 36, 64, 100]

// Array slice with range
let letters = ["a", "b", "c", "d", "e"];
let mid = letters[1..4];    // ["b", "c", "d"]
print(mid);
```

Shebang — executable scripts

A `.rhai` file can be made directly executable by adding a shebang as the first line:

```
#!/usr/bin/env -S recon --script
```

The `-S` flag instructs `/usr/bin/env` to split its argument on whitespace, so `recon` receives `--script` as a separate flag. This works on macOS and all major Linux distributions.

Full example

```
cat > ~/bin/health <<'EOF'
#!/usr/bin/env -S recon --script
let host = if args.len() > 1 { args[1] } else { "example.com" };
let r = https(`https://${host}`);
print(`${r.status} ${host} (${r.duration_ms}ms)`);
return if r.status == 200 { 0 } else { 1 };
EOF
chmod +x ~/bin/health

./health example.com          # run directly
./health api.example.com      # host from args[1]
```

How it works

When the kernel executes a shebang file, it calls the interpreter with the script path as an argument — equivalent to `recon --script ./health`. Recon detects a leading `#!` in the source and replaces it with `//` (a Rhai line comment) before compilation. This preserves line numbers in error messages while making the file valid Rhai.

Arguments and flags

Trailing arguments after the script name land in `args[1..]` exactly as with `--script`. CLI flags (`-k`, `-v`, `--connect-timeout`, etc.) set before or after the script name in the invocation are reflected in the `flags` map inside the script.

```
./health api.example.com          # args[1] = "api.example.com"
recon -k ./health api.example.com # flags.insecure = true inside script
```

Notes

- On older systems where `/usr/bin/env -S` is not available, use the full path:
`#!/usr/local/bin/recon --script` (hard-coded path, no `-S` needed).
- The shebang is only stripped when it appears on the **first line**. A `#!` appearing anywhere else in the file remains a parse error (Rhai does not treat `#` as a comment character).
- `recon --help shebang` summarises the shebang feature.

CLI inheritance

When a script runs, two read-only constants land at the top of its scope:

- **args** — an array of strings. `args[0]` is the script path; `args[1..]` are the positional arguments after `--script PATH`.
- **flags** — a map mirroring the relevant CLI flags (lower-case, snake_case keys). Useful for scripts that want to respect `-v`, `-k`, `--connect-timeout`, etc. at invocation time.

```
// Reasonable defaults, overridable via args
let url = if args.len() > 1 { args[1] } else { "https://example.com/" };
let timeout = flags.connect_timeout; // from -C / --connect-timeout

let r = http(url, #{ timeout_ms: timeout * 1000 });
print(`${r.status} ${r.url}`);
```

Bindings that take their own opts map (e.g. `http(url, opts)`) layer the caller's opts ON TOP of the CLI defaults. If the caller wants to ignore the CLI's `-H` header list, they can pass `headers: []` in their opts.

Core helpers

Exposed by `src/script/bindings/helpers.rs`. Always available.

Function	Returns	Description
<code>sleep_ms(ms)</code>	—	Block the current thread.
<code>sleep(ms)</code>	—	Alias of <code>sleep_ms</code> (added in 0.56.0 for thread-side readability).
<code>env(name)</code>	string	Process environment variable; <code>""</code> if unset.
<code>env(name, default)</code>	string	With fallback.
<code>env_all()</code>	Map	Snapshot every process env var. Aliased as <code>envAll</code> .
<code>load_dotenv(path)</code>	int	Parse <code>.env</code> and set each KEY=VALUE in the process env (overrides existing). Returns count of vars set. Aliased as <code>loadDotEnv</code> .
<code>load_dotenv(path, override)</code>	int	Two-arg form: <code>false</code> leaves pre-existing env vars in place.
<code>script_path</code>	string (constant)	Resolved absolute path of the running script. Pushed into the Scope at startup.
<code>script_dir</code>	string (constant)	Parent directory of <code>script_path</code> . Use to load sibling files independent of CWD: <code>load_dotenv(script_dir + "/.env")</code> .
<code>script_name</code>	string (constant)	File stem of the running script (basename minus extension). Useful with the per-script overlay convention: <code>load_dotenv(script_dir + "/.env." + script_name)</code> .
<code>now()</code>	int	Unix seconds.
<code>now_ms()</code>	int	Unix milliseconds.
<code>assert(cond, msg)</code>	—	Throws on false.
<code>json_parse(s)</code>	Dynamic	Parse JSON text → native Rhai value.
<code>json_stringify(v)</code>	string	Serialize to compact JSON.
<code>json_stringify(v, pretty)</code>	string	<code>pretty: true</code> = 2-space indent.
<code>json_stringify(v, n)</code>	string	<code>n</code> = spaces per indent (clamped 1..=8).

Rhai's built-in `print(x)` writes `x` + newline via the engine's debug callback. `recon` adds byte-precise siblings (see next section).

Examples

```

// Probe a URL, fail fast
let r = http("https://api.example.com/health");
assert(r.status == 200, `expected 200, got ${r.status}`);

// JSON round-trip
let obj = json_parse(r.body);
print(json_stringify(obj, true));    // pretty

// Time gate
let t0 = now_ms();
let r = http("https://slow.example.com/");
let dt = now_ms() - t0;
print(`took ${dt} ms`);

// Polling with timeout
let deadline = now() + 30;
while now() < deadline {
    let r = http("https://example.com/ready");
    if r.status == 200 { return 0; }
    sleep_ms(500);
}
return 1;

// Layered .env loading: common defaults, per-script overrides.
// Use script_dir so the .env files don't depend on CWD – they live
// alongside the script in the same directory.
load_dotenv(script_dir + "/.env");           // shared
load_dotenv(script_dir + "/.env." + script_name); // per-script
overlay wins
print(env("LOG_LEVEL"));

// Non-override variant: shell-env wins over file values
load_dotenv(script_dir + "/.env", false);

// env_all() – useful for filtering or logging the entire env
let everything = env_all();
print(`process env has ${everything.len()} variables`);

```

`load_dotenv` overrides existing env values by default — that's what makes the `common.env`, then `.env.<scriptname>` pattern layer correctly. Pass `false` as a second arg to leave pre-existing values in place (e.g. when shell exports should take priority over file defaults). `std::env::set_var` is technically unsound under concurrent reads on some platforms, so call `load_dotenv` at the top of the script — before `thread_spawn` or any concurrent binding.

Output bindings

0.53.0. Byte-precise output that Rhai's built-in `print()` can't do.

Function	Description
<code>print_raw(s) / print_raw(blob)</code>	Write to stdout without a trailing newline; flush.
<code>eprint(s)</code>	Write to stderr with a newline.
<code>eprint_raw(s) / eprint_raw(blob)</code>	Write to stderr without newline; flush.
<code>flush()</code>	Explicit stdout flush.

Examples

```
// Progress indicator
print_raw("loading");
for i in 0..10 {
    sleep_ms(200);
    print_raw(".");
}
print_raw("\n");

// Send bytes to stdout (e.g. for piping into another tool)
let bytes = http("https://example.com/image.png").body;
print_raw(bytes);

// Log to stderr; keep stdout for data
eprint(`[${now_ms()}] starting probe...`);
print_raw(json_stringify(result));      // stdout stays parseable
```

File I/O bindings

0.53.0. Both whole-file convenience and a streaming handle API.

Whole-file helpers

Function	Description
<code>file_read(path)</code>	Read entire file → Blob.
<code>file_write_all(path, blob str)</code>	Overwrite. Returns bytes written.
<code>file_append_all(path, blob str)</code>	Append; creates if absent.
<code>file_exists(path)</code>	Bool.
<code>file_size(path)</code>	Bytes.
<code>file_delete(path)</code>	Unlink.

`path` accepts a filesystem path or a `file://` URL.

Streaming handle API

Function	Description
<code>file_open(path, mode)</code>	Returns a <code>FileHandle</code> . Modes: <code>r</code> , <code>w</code> (truncate+create), <code>rw</code> , <code>rw+</code> / <code>w+</code> (read+write+create+truncate), <code>a</code> (append+create), <code>ra</code> (append+read).
<code>file_read(h, n)</code>	Read up to <code>n</code> bytes → Blob.
<code>file_read_all(h)</code>	Drain to Blob.
<code>file_write(h, blob str)</code>	Write all bytes. Returns count.
<code>file_seek(h, pos, whence)</code>	<code>whence</code> = <code>start</code> <code>cur</code> <code>end</code> . Returns new position.
<code>file_tell(h)</code>	Current position.
<code>file_flush(h)</code>	Sync buffered writes.
<code>file_close(h)</code>	Close (drops the underlying file).

Examples

```

// Whole-file
let data = file_read("config.json");
file_write_all("/tmp/copy.json", data);
file_append_all("/tmp/audit.log", `probe at ${now()}\n`);

// Check before read
if !file_exists("secrets.key") {
    eprint("missing key");
    return 2;
}

// Streaming – copy with 4KB chunks
let src = file_open("big.bin", "r");
let dst = file_open("/tmp/big.bin.copy", "w");
loop {
    let chunk = file_read(src, 4096);
    if chunk.len() == 0 { break; }
    file_write(dst, chunk);
}
file_close(src);
file_close(dst);

// Random access
let h = file_open("/tmp/log.bin", "rwc");
file_write(h, "record-1\n");
file_write(h, "record-2\n");
file_seek(h, 0, "start");
let head = file_read_all(h);
print(`contents: ${head.len()} bytes`);
file_close(h);

```

Clipboard bindings

0.70.0. The `clipboard` static module exposes the system clipboard for read/write access from Rhai scripts. Backed by the same `arboard` crate that powers the `--clipboard` family of CLI flags, so behaviour is identical across platforms (macOS pasteboard, Linux X11/Wayland, Windows).

Function	Returns	Description
<code>clipboard::get()</code>	string	Current clipboard text. Errors if no clipboard available or content isn't text.
<code>clipboard::set(text)</code>	()	Replace clipboard contents with <code>text</code> .

Examples

```
// Read, transform, write back to clipboard.
let original = clipboard::get();
let upper = original.to_upper();
clipboard::set(upper);
print("Replaced clipboard with uppercase form");
```

```
// Fetch a URL and place the prettified JSON response on the clipboard.
let r = http("https://api.example.com/data");
let pretty = json_stringify(json_parse(r.body.to_string()));
clipboard::set(pretty);
print("Result copied to clipboard");
```

Comparison bindings

0.53.0. Diff two in-memory values (strings or Blobs).

Function	Returns	Description
<code>compare(a, b)</code>	Map	<code>{ identical, binary, a_bytes, b_bytes, added, removed, diff }</code> .

Both `a` and `b` accept strings OR Blobs. The binding does NOT fetch URLs — scripts should pre-load via `http()` or `file_read()` and pass the bytes.

Examples


```
// Compare two API responses
let a = http("https://api.v1.example.com/users/42").body;
let b = http("https://api.v2.example.com/users/42").body;
let r = compare(a, b);
if r.identical {
    print("parity OK");
    return 0;
}
print(`diverged: +${r.added} / -${r.removed}`);
print_raw(r.diff);
return 1;

// Compare a live URL against a baseline file
let live = http("https://example.com/status").body;
let baseline = file_read("status.baseline").to_string();
let r = compare(live.to_string(), baseline);
if !r.identical { print_raw(r.diff); }

// Binary content
let old = file_read("logo.v1.png");
let new = file_read("logo.v2.png");
let r = compare(old, new);
print(`binary: ${r.binary}, size delta: ${r.b_bytes - r.a_bytes}`);
```

HTTP binding

The workhorse. Two call forms:

```
http(url)           → response map
http(url, opts)     → response map
```

Response map

Key	Type	Description
status	int	HTTP status code.
url	string	Original URL.
final_url	string	After redirects.
headers	Map	Response headers; values are strings (first) or arrays when repeated.
body	Blob string	Response body. String when content-type is text-ish; Blob otherwise.

Key	Type	Description
<code>duration_ms</code>	int	Total wall-clock time.
<code>http_version</code>	string	HTTP/1.1 , HTTP/2 , etc.

Options map

Every key is optional.

Key	Type	Description
<code>method</code>	string	GET, POST, PUT, PATCH, DELETE, HEAD.
<code>headers</code>	Map	Key → value (or array of values). Merged with CLI <code>-H</code> defaults.
<code>body</code>	string / Blob	Request body.
<code>json</code>	Dynamic	Serialize as JSON, set <code>Content-Type</code> + <code>Accept</code> .
<code>form</code>	Map	URL-encoded form body.
<code>multipart</code>	Array of Maps	Each Map has <code>name</code> , <code>value</code> or <code>file</code> , optional <code>filename</code> , <code>content_type</code> .
<code>basic_auth</code>	string	<code>"user:pass"</code> .
<code>bearer</code>	string	Bearer token.
<code>user_agent</code>	string	Override User-Agent.
<code>referer</code>	string	Referer.
<code>timeout_ms</code>	int	Total operation timeout.
<code>connect_timeout</code>	int	TCP connect timeout (seconds).
<code>insecure</code>	bool	Skip TLS verification.
<code>follow_redirects</code>	bool	Enable / disable redirect follow.
<code>max_redirects</code>	int	Redirect cap.
<code>compressed</code>	bool	Request + auto-decompress.
<code>cookiejar</code>	string	Path to a SQLite cookie jar.
<code>proxy</code>	string	Proxy URL.
<code>proxy_user</code>	string	Proxy basic auth.
<code>noproxy</code>	string	Bypass list.
<code>proxy_insecure</code>	bool	Skip cert verification on proxy.
<code>proxy_cacert</code>	string	Extra CA for proxy.

Key	Type	Description
<code>unix_socket</code>	string	Route via Unix socket.
<code>hsts</code>	string	Path to HSTS cache.
<code>client_cert</code>	string	PEM client cert.
<code>client_key</code>	string	PEM client key.
<code>cert_type / key_type</code>	string	PEM DER ENG (see CLI docs).
<code>pass</code>	string	PKCS#8 passphrase (reserved).
<code>wait</code>	int	(0.67.0) Seconds between URLs in batch mode (CLI-only effect).
<code>tries</code>	int	(0.67.0) Total attempts; overrides <code>retry</code> . <code>tries = retries + 1</code> .
<code>accept</code>	string	(0.67.0) Comma-separated filename-suffix accept list.
<code>reject</code>	string	(0.67.0) Comma-separated filename-suffix reject list.
<code>prettify</code>	bool	Pretty-print response body (auto-detect format).
<code>prettify_as</code>	string	Force prettify format (json/xml/html/yaml/csv/tsv/auto). Implies <code>prettify: true</code> .
<code>render</code>	bool	(0.96.0) Render <code>text/html</code> response bodies to text; non-HTML passes through. <code>resp.body</code> becomes the rendered text.
<code>width</code>	int	(0.96.0) Wrap column for rendered output (default: terminal width on a TTY, else 80).
<code>render_color</code>	string	(0.96.0) ANSI styling of rendered text: "auto" / "always" / "never" (default "never" in scripts).
<code>impersonate</code>	string	(0.77.0) Browser TLS+H2 fingerprint profile name (e.g. "chrome_131", "firefox_128", "safari_17.5"). Requires a build with <code>--features impersonate</code> ; rejected with a rebuild hint otherwise. Hyphens accepted ("chrome-131" $\equiv$ "chrome_131"). See <code>recon --help impersonate</code> .
<code>ja3</code>	string	(0.77.0, deferred) Reserved for raw JA3 ClientHello override. Errors at runtime as not-yet-implemented; use <code>impersonate</code> instead.
<code>ja4</code>	string	(0.77.0, deferred) Reserved for raw JA4 fingerprint override. Errors at runtime as not-yet-implemented.

Key	Type	Description
<code>http2_fingerprint</code>	string	(0.77.0, deferred) Reserved for raw Akamai HTTP/2 fingerprint override. Errors at runtime as not-yet-implemented.

Examples

```

// Simple GET
let r = http("https://api.example.com/status");
print(`${r.status} ${r.body.len()} bytes`);

// POST JSON
let r = http("https://api.example.com/users", #{
  method: "POST",
  json: #{ name: "alice", email: "a@example.com" },
});
assert(r.status == 201, "create failed");

// Form POST
http("https://api.example.com/login", #{
  method: "POST",
  form: #{ user: "alice", pass: "s3cr3t" },
  cookiejar: "/tmp/jar.db",
});

// Multipart upload
http("https://upload.example.com/avatars", #{
  method: "POST",
  multipart: [
    #{ name: "title", value: "profile" },
    #{ name: "image", file: "avatar.png", content_type: "image/png" },
  ],
  bearer: env("API_TOKEN"),
});

// Bearer auth + timeout + compressed
let r = http("https://api.example.com/items", #{
  bearer: env("API_TOKEN"),
  timeout_ms: 5000,
  compressed: true,
  headers: #{ "Accept": "application/json" },
});

// Follow + retry on 5xx
for attempt in 0..3 {
  let r = http("https://api.example.com/", #{ follow_redirects: true });
  if r.status < 500 { return r.status == 200 ? 0 : 1; }
  sleep_ms(1000 * (attempt + 1));
}
return 2;

// Proxy + HSTS
http("http://example.com/", #{
  proxy: "socks5h://127.0.0.1:9050",
  hsts: "/tmp/hsts.txt",

```

```
});

// Browser fingerprint impersonation (0.77.0; requires --features
impersonate)
let r = http("https://tls.peet.ws/api/all", #{
    impersonate: "chrome_131",
});
print(`status: ${r.status}, body: ${r.body.len()} bytes`);

// mTLS
http("https://mtls.example.com/", #{
    client_cert: "/etc/keys/bundle.pem",
});

// Unix socket
let r = http("http://localhost/v1.40/version", #{
    unix_socket: "/var/run/docker.sock",
});
print(json_stringify(json_parse(r.body), 2));

// Pretty-print response body (force JSON format)
let r = http("https://api.example.com/data", #{ prettify_as: "json" });
print(r.body);
```

Browser (sticky-session) binding

A stateful HTTP "browser" — cookies and headers persist across calls against the same handle. Matches the common "log in once, then make authenticated requests" pattern.

Function	Description
<code>browser()</code>	New browser with a tempfile cookie jar.
<code>browser(opts)</code>	With initial config.
<code>use_persistent_session(name)</code>	Browser method; jar is <code>~/.recon/jars/<name>.db</code> .
<code>
.get(url [, opts])</code>	GET via the browser.
<code>
.post(url, opts)</code>	POST.
<code>
.put(url, opts) / .patch(...) / .delete(url [, opts])</code>	Other methods.
<code>
.request(method, url, opts)</code>	Catch-all.
<code>
.header(name, value)</code>	Sticky request header.

Function	Description
<code>
.headers(map)</code>	Merge sticky headers.
<code>
.clear_headers()</code>	Remove sticky headers.
<code>
.user_agent(str)</code>	Override User-Agent for the lifetime of the browser.
<code>
.basic_auth(user, pass)</code>	Sticky Basic auth.
<code>
.timeout_ms(n) / .connect_timeout(secs)</code>	Sticky timeouts.
<code>
.insecure(bool) / .follow_redirects(bool) / .max_redirects(n)</code>	Sticky knobs.
<code>
.cookiejar()</code>	Current jar path.
<code>
.cookies()</code>	Array of cookie records from the jar.
<code>
.cookie_set(name, value, domain, path)</code>	Inject a cookie.

Examples

```
// Login, then fetch protected resources
let br = browser();
br.user_agent("recon-test/0.58");
br.post("https://example.com/login", #{
  form: #{ user: "alice", pass: "s3cr3t" },
});
let r = br.get("https://example.com/dashboard");
assert(r.status == 200, "dashboard fetch failed");

// Persist across runs – next invocation sees the same cookies
let br = browser();
br.use_persistent_session("gh-session");
br.header("Accept", "application/vnd.github+json");
let r = br.get("https://api.github.com/user");

// Multi-persona fan-out
let personas = [{ name: "a", jar: "persona-a" }, { name: "b", jar:
"persona-b" }];
for p in personas {
  let br = browser();
  br.use_persistent_session(p.jar);
  let r = br.get("https://api.example.com/whoami");
  print(`${p.name}: ${r.status}`);
}
```

TCP, TCP-server, UDP

0.57.0 shipped the server bindings. Handles are Send+Sync so they survive `thread_spawn` (0.56.0) boundaries — the expected pattern is "accept on main, spawn a handler per connection".

TCP client (existing)

Function	Description
<code>tcp(url)</code>	TCP connect probe; returns <code>#{ status, duration_ms, ... }</code> .
<code>tcp(url, opts)</code>	With <code>timeout</code> (ms).

TCP server (0.57.0)

Function	Description
<code>tcp_listen(addr)</code>	Bind. <code>addr</code> like <code>"0.0.0.0:8080"</code> or <code>"[::]:8080"</code> .

Function	Description
<code>tcp_accept(listener)</code>	Blocking accept. Returns a <code>TcpConn</code> .
<code>tcp_accept(listener, timeout_ms)</code>	Times out with an error.
<code>tcp_read(conn, n, timeout_ms)</code>	Read up to N bytes → <code>Blob</code> .
<code>tcp_read_line(conn, timeout_ms)</code>	<code>\n</code> -terminated line; CR/LF stripped.
<code>tcp_write(conn, blob str)</code>	Write all. Returns bytes.
<code>tcp_peer_addr(conn)</code>	Remote <code>SocketAddr</code> string.
<code>tcp_close(conn)</code>	Close connection.
<code>tcp_close_listener(l)</code>	Close listener.

UDP

Function	Description
<code>udp_bind(addr)</code>	Bind.
<code>udp_recv_from(sock, max_len)</code>	Block until a datagram arrives. Returns <code>#{ data: Blob, addr: string }</code> .
<code>udp_recv_from(sock, max_len, timeout_ms)</code>	With timeout.
<code>udp_send_to(sock, blob str, addr)</code>	Returns bytes sent.
<code>udp_close(sock)</code>	Release.

Examples

```
// TCP client probe
let r = tcp("example.com:443");
print(`${r.status} in ${r.duration_ms}ms`);

// Concurrent TCP echo server (from script/tcp-echo.rhai)
let l = tcp_listen("127.0.0.1:9000");
loop {
    let conn = tcp_accept(l);
    thread_spawn(|c| {
        let peer = tcp_peer_addr(c);
        let line = tcp_read_line(c, 5000);
        tcp_write(c, `echo: ${line}` + "\n");
        tcp_close(c);
    }, conn);
}

// UDP listener (from script/udp-listen.rhai)
let s = udp_bind("127.0.0.1:9001");
loop {
    let r = udp_recv_from(s, 65536, 30000);
    print(`${r.addr} → ${r.data.len()} bytes`);
}

// UDP beacon sender
let s = udp_bind("0.0.0.0:0");
for i in 0..5 {
    udp_send_to(s, `beacon-${i}`, "127.0.0.1:9001");
    sleep_ms(200);
}
udp_close(s);
```

Threading bindings

0.56.0. Requires rhai's `sync` feature (enabled). `spawn` alone is reserved by Rhai — use `thread_spawn`.

Function	Description
<code>thread_spawn(fn_ptr)</code>	Spawn a closure. Returns a <code>ThreadHandle</code> .
<code>thread_spawn(fn_ptr, arg)</code>	With one forwarded arg.
<code>thread_spawn(fn_ptr, args_array)</code>	With N args.
<code>join(h)</code>	Block; returns the closure's return value (or the worker's error).

Function	Description
<code>tid()</code>	Current thread id (stable within a run).
<code>sleep(ms)</code>	Alias of <code>sleep_ms</code> .
<code>channel()</code>	Unbounded MPSC — returns <code>[sender, receiver]</code> .
<code>channel_bounded(n)</code>	Bounded; <code>try_send</code> returns false when full.
<code>send(tx, val)</code>	Blocking send.
<code>try_send(tx, val)</code>	Non-blocking; false on full.
<code>recv(rx)</code>	Blocking.
<code>recv(rx, timeout_ms)</code>	With timeout.
<code>try_recv(rx)</code>	Non-blocking; <code>()</code> when empty.

Examples

```

// Fan out probes, gather via channel
let c = channel();
let tx = c[0];
let rx = c[1];

let urls = [
    "https://a.example.com/",
    "https://b.example.com/",
    "https://c.example.com/",
];

for url in urls {
    thread_spawn(|u| {
        let r = http(u, #{ timeout_ms: 5000 });
        send(tx, `${u} → ${r.status}`);
    }, url);
}

for i in 0..urls.len() {
    print(recv(rx, 10000));
}

// Bounded channel – producer/consumer with back-pressure
let c = channel_bounded(4);
let tx = c[0];
let rx = c[1];

thread_spawn(|| {
    for i in 0..100 { send(tx, i); }
});

let sum = 0;
for j in 0..100 {
    sum += recv(rx, 1000);
}
print(`total: ${sum}`);

// Join for return value
let h = thread_spawn(|| {
    // Do something heavy
    let r = http("https://slow.example.com/");
    r.status
});
let status = join(h);
print(`completed with status ${status}`);

```

Shell binding

Run external commands from a script. Two forms — `shell()` blocks and captures everything, `shell_stream()` fires a callback per line as the child writes it. Built for run-and-watch scripts (local software updates, multi-step setup, log filtering) and as the substrate for the upcoming TUI pane primitive.

Function	Description
<code>shell(cmd_string) / shell(cmd_string, opts)</code>	Run via <code>sh -c <s></code> (Unix) or <code>cmd /C <s></code> (Windows). Pipes / globs / redirects / <code>&&</code> chains work.
<code>shell(argv_array) / shell(argv_array, opts)</code>	Direct argv form. No shell layer, no quoting surprises.
<code>shell_stream(cmd, callback) / shell_stream(cmd, callback, opts)</code>	Streaming form. Callback fires per merged stdout+stderr line as the child writes; returns the exit code on child exit.

Return value (`shell`)

A Map:

Key	Type	Description
<code>stdout</code>	String	Captured stdout (lossy UTF-8).
<code>stderr</code>	String	Captured stderr (or empty when <code>merge_stderr: true</code>).
<code>exit_code</code>	Int	Child's exit code, or <code>-1</code> if killed by signal.
<code>success</code>	Bool	<code>exit_code == 0</code> .

Opts map (both forms)

All keys optional.

Key	Type	Behaviour
<code>cwd</code>	String	Working directory (default: inherit).
<code>env</code>	Map	Name → value, layered on top of the parent environment.
<code>env_clear</code>	Bool	Drop the parent env entirely before applying <code>env</code> .
<code>timeout_ms</code>	Int	Kill the child after N ms; raises a Rhai error the script can <code>try / catch</code> .
<code>merge_stderr</code>	Bool	<code>shell</code> only — fold stderr into stdout. <code>shell_stream</code> always merges.

Examples

```
// Blocking – common case.
let r = shell("git status --short");
if r.success { print(r.stdout); } else { print(`exit ${r.exit_code}:
${r.stderr}`); }

// argv form skips shell expansion.
let r = shell(["echo", "$HOME"]); // stdout: "$HOME\n"

// Opts.
let r = shell("cargo test", #{
  cwd: "/path/to/repo",
  env: #{ RUST_LOG: "info" },
  timeout_ms: 60000,
});

// Streaming – callback per line as the child writes.
let exit = shell_stream("brew upgrade", |line| {
  print(`[brew] ${line}`);
});
print(`brew exit ${exit}`);

// Multi-step update with error isolation.
for cmd in ["brew upgrade", "npm -g update", "cargo install-update -a"] {
  try {
    shell_stream(cmd, |line| print(line));
  } catch (e) {
    print(`step '${cmd}' failed: ${e}`);
  }
}
```

A non-zero exit code does NOT raise an error — check `r.success` (or compare `exit_code` in the streaming form's return value). A timeout DOES raise an error, hence the `try / catch` in the multi-step example.

TUI dashboard binding

Multi-pane text dashboard for scripts. Sits on top of `shell_stream`: the canonical use case is a run-and-watch script where one or more subprocesses stream their output into distinct panes while the script logs its own progress.

Function	Description
<code>tui::run(callback)</code>	Enter alt-screen, build a Dashboard, pass it to the callback. Restores the terminal on any exit (normal, Rhai

Function	Description
	error, panic, SIGINT). Errors if stdout isn't a TTY or another dashboard is already active.
<code>d.split_vertical([p1, p2, ...])</code>	Stack panes top-to-bottom. Each percentage is in <code>(0, 100]</code> . Returns an Array of pane handles. Can only be called once per dashboard.
<code>d.split_horizontal([p1, p2, ...])</code>	Lay panes left-to-right. Same semantics as <code>split_vertical</code> .
<code>pane.println(line)</code>	Append a line to the pane's scrollbar. Auto-scrolls. Cap: 1000 lines per pane.
<code>pane.title(s)</code>	Set the pane's top-row title. Rendered as <code><s></code> followed by a horizontal rule.
<code>pane.clear()</code>	Empty the pane's scrollbar. Title is preserved.

Example

```
tui::run(|d| {
  let parts = d.split_vertical([70, 30]);
  let main = parts[0];
  let status = parts[1];

  main.title("subprocess output");
  status.title("script progress");

  for cmd in ["brew upgrade", "npm -g update", "cargo install-update -a"]
  {
    status.println(`[running] ${cmd}`);
    try {
      shell_stream(cmd, |line| main.println(line));
      status.println(`[done] ${cmd}`);
    } catch (e) {
      status.println(`[failed] ${cmd}: ${e}`);
    }
  }
  status.println("all done");
  sleep_ms(1500);
});
```

v1 limitations (intentional)

- **No raw mode.** Terminal resize is handled by a 150 ms polling redraw rather than via SIGWINCH events. Layout glitches between updates are possible until the next pane write.

- **No PTY allocation.** Subprocesses spawned via `shell_stream` still detect "not a terminal" and switch to non-coloured / unbuffered output. Fixing this requires a real PTY (`portable-pty` / `nix`) and is its own work item.
- **Truncation, not wrapping.** Lines longer than the pane width are cut at the right edge.
- **Wide characters undercounted.** Emoji and CJK render as two terminal cells but count as one when truncating. Lines containing them may overflow by a column or two.

DNS binding

Function	Description
<code>dns(host)</code>	Default bundle: A, AAAA, MX, TXT, NS, CNAME, SOA. Returns a Map keyed by record type.
<code>dns(host, types_str)</code>	Custom types: "A,AAAA,MX" .
<code>dns(host, types_str, opts)</code>	<code>opts</code> may contain <code>servers: "1.1.1.1,8.8.8.8"</code> .

Examples

```
let r = dns("example.com");
print(`A: ${r.A.join(", ")}`);
print(`MX: ${r.MX.join(", ")}`);

// Custom resolver
let r = dns("example.com", "A,AAAA", #{ servers: "127.0.0.1:5353" });
print(r);

// Check DNSSEC chain presence
let r = dns("example.com", "DNSKEY,DS");
if r.DNSKEY.len() == 0 { print("no DNSKEY"); }
```

TLS probe binding

Function	Description
<code>tls(host [, opts])</code>	Handshake + cert introspection. Returns a Map with <code>subject</code> , <code>issuer</code> , <code>sans</code> , <code>not_before</code> , <code>not_after</code> , <code>days_remaining</code> , <code>fingerprints</code> , <code>chain</code> , <code>tls_version</code> , <code>cipher_suite</code> .

Examples


```

let t = tls("example.com:443");
print(`CN: ${t.subject}`);
print(`issuer: ${t.issuer}`);
print(`expires in ${t.days_remaining} days`);
if t.days_remaining < 14 { eprint("cert expiring soon"); return 1; }

// Verify an expected subject
let t = tls("api.example.com:443");
assert(t.sans.contains("api.example.com"), "SAN mismatch");

// Per-host SNI override
let t = tls("alt.example.com:443", #{ sni: "primary.example.com" });

```

Ping binding

Function	Description
<code>ping(host)</code>	Default TCP ping, 4 packets. Returns <code>#{ sent, received, loss_pct, min_ms, avg_ms, max_ms, ...}</code> .
<code>ping(host, opts)</code>	<code>opts</code> supports <code>count</code> , <code>icmp: true</code> , <code>timeout_ms</code> .

Examples

```

let r = ping("8.8.8.8");
print(`${r.received}/${r.sent} received, avg ${r.avg_ms}ms`);

// ICMP (needs privileges on macOS/Linux for non-DGRAM types)
let r = ping("example.com", #{ icmp: true, count: 10 });

// TCP ping to a specific port
let r = ping("example.com:443");
assert(r.received > 0, "no reachability");

```

File transfer bindings

FTP

Function	Description
<code>ftp(url)</code>	FTP / FTPS anonymous listing or retrieve.
<code>ftp(url, opts)</code>	<code>opts</code> : <code>user</code> , <code>pass</code> , <code>ftps_implicit</code> .

SFTP

Function	Description
<code>sftp(url)</code>	SSH-backed listing or retrieve.
<code>sftp(url, opts)</code>	<code>opts</code> : <code>user</code> , <code>pass</code> , <code>privkey</code> , <code>port</code> .

TFTP

Function	Description
<code>tftp(url)</code>	RFC 1350 download.
<code>tftp(url, opts)</code>	<code>opts</code> : <code>blksize</code> (RFC 2348).

Gopher

Function	Description
<code>gopher(url)</code>	RFC 1436 selector fetch.

Examples

```
let listing = ftp("ftp://ftp.example.com/");
print(`${listing.entries.len()} entries`);

let data = ftp("ftp://ftp.example.com/file.bin");
print(`${data.size} bytes`);

let h = sftp("sftp://me@example.com/srv/app.log", #{ privkey:
"~/.ssh/ci_rsa" });
file_write_all("/tmp/app.log", h.body);

let img = tftp("tftp://tftp.example.com/boot.img", #{ blksize: 8192 });
print(img.size);

let g = gopher("gopher://gopher.floodgap.com/0/gopher/proxy");
print_raw(g.body);
```

Mail retrieval bindings

POP3

Function	Description
<code>pop3(url)</code>	Capability + message listing.
<code>pop3(url, opts)</code>	<code>opts: stls</code> , <code>retrieve_index</code> (int, 1-based).

IMAP

Function	Description
<code>imap(url)</code>	EXAMINE + capabilities.
<code>imap(url, opts)</code>	<code>opts: fetch</code> (int — message index), <code>peek: true</code> .

Examples

```
// POP3 probe with STLS
let r = pop3("pop3://alice:s3cr3t@pop.example.com/", #{ stls: true });
print(`${r.message_count} messages`);

let first = pop3s("pop3s://alice:s3cr3t@pop.example.com/", #{
  retrieve_index: 1 });
print_raw(first.body);

// IMAP peek
let r = imap("imaps://alice:s3cr3t@imap.example.com/INBOX", #{ fetch: 3,
  peek: true });
print_raw(r.body);
```

SMTP binding

Function	Description
<code>smtp(url)</code>	Capability + STARTTLS probe.
<code>smtp(url, opts)</code>	Send a message. <code>opts</code> keys: <code>from</code> , <code>to</code> (array), <code>subject</code> , <code>body</code> , <code>headers</code> , <code>auth</code> ("user:pass"), <code>helo</code> , <code>no_starttls</code> , <code>dkim_key</code> , <code>dkim_selector</code> , <code>dkim_domain</code> .

Examples

```
// Probe only
let r = smtp("smtp://mail.example.com:25");
print(`capabilities: ${r.capabilities.join(", ")}`);

// Send
smtp("smtp://smtp.example.com:587", #{
  from: "me@example.com",
  to: ["you@example.com"],
  subject: "Automated notice",
  body: "Probe result: OK",
  auth: env("SMTP_CREDS"),
});

// DKIM-signed
smtp("smtp://smtp.example.com:587", #{
  from: "me@example.com",
  to: ["you@example.com"],
  subject: "Signed",
  body: "This message is DKIM-signed.",
  dkim_key: "~/dkim/default.pem",
  dkim_selector: "default",
});
```

WebSocket binding

Function	Description
<code>ws(url)</code>	Handshake + Ping/Pong. Returns <code>#{ status, duration_ms, negotiated_protocol }</code> .
<code>ws(url, opts)</code>	<code>opts: headers, subprotocols, send_text, send_binary.</code>

Examples

```
let r = ws("wss://echo.websocket.events");
print(`connected in ${r.duration_ms}ms`);

// Send a frame, check response
let r = ws("wss://echo.websocket.events", #{ send_text: "hello" });
print(r.received);
```

Other protocol bindings

Brief interface reference; each has at least one example in `script/*.rhai`.

LDAP

```
let r = ldap("ldap://ldap.example.com");           // anonymous RootDSE query
```

RTSP

```
let r = rtsp("rtsp://camera.example.com/stream1");
```

Dict

```
let r = dict("dict://dict.org/d:recon");
```

NTP

```
let r = ntp("ntp://pool.ntp.org");  
print(`offset: ${r.offset_ms}ms, rtt: ${r.rtt_ms}ms`);
```

Memcached

```
let r = memcached("memcache://cache.example.com/?version");  
let s = memcached("memcache://cache.example.com/", #{ command: "stats" });
```

Redis

```
let r = redis("redis://cache.example.com/");       // PING  
let info = redis("redis://cache.example.com/", #{ command: "INFO server"  
});
```

MQTT

```
let r = mqtt("mqtt://broker.example.com", #{  
    topic: "sensors/room-1",  
    message: '{"temp": 21.3}',  
});
```

Encoding & decoding bindings

0.55.0 expanded encode with Aztec + PDF417 and added decode.

Function	Description
<code>encode::qr(text)</code>	PNG Blob.
<code>encode::datamatrix(text)</code>	PNG Blob.
<code>encode::barcode(format, text)</code>	PNG Blob for any 1D/2D.
<code>encode::encode(format, text)</code>	PNG Blob (default renderer).
<code>encode::encode(format, text, output_format)</code>	<code>output_format = ascii / svg / png</code> . Returns string for ascii/svg, Blob for png.
<code>encode::encode(format, text, output_format, opts)</code>	<code>opts: qr_level: "L" "M" "Q" "H"</code> .
<code>encode::decode(blob)</code>	Scan an image Blob → <code>{ text, format }</code> .
<code>encode::list()</code>	Supported format names.

Examples

```
// Generate QR
let png = encode::qr("https://example.com");
file_write_all("/tmp/site.png", png);

// Tune QR durability
let png = encode::encode("qr", "sticker-text", "png", #{ qr_level: "H" });

// Get SVG
let svg = encode::encode("qr", "recon", "svg");
file_write_all("/tmp/site.svg", svg);

// Decode
let data = file_read("/tmp/site.png");
let r = encode::decode(data);
print(`${r.format}: ${r.text}`);

// Round-trip test
let png = encode::qr("hello");
let r = encode::decode(png);
assert(r.text == "hello", "round-trip failed");
```

Hashing binding

Ten algorithms plus CRC32.

Function	Description
<code>md5(data)</code>	MD5 (hex string).
<code>sha1(data)</code>	SHA-1.
<code>sha224</code> / <code>sha256</code> / <code>sha384</code> / <code>sha512</code>	SHA-2 family.
<code>sha3_256</code> / <code>sha3_512</code>	SHA-3.
<code>blake3(data)</code>	BLAKE3.
<code>crc32(data)</code>	CRC32 (hex, 8 chars).
<code>hmac_sha256(key, data)</code>	Keyed HMAC.

`data` accepts string or Blob.

Examples

```
let h = sha256(file_read("big.iso"));
print(h);

// HMAC
let sig = hmac_sha256("sharedsecret", "payload");
print(sig);

// Hash an HTTP body
let r = http("https://example.com/");
let fp = blake3(r.body);
print(`fingerprint: ${fp}`);
```

Compression & archive bindings

Compression

Function	Description
<code>compression::compress(algo, data [, level])</code>	<code>algo</code> = <code>gzip</code> , <code>deflate</code> , <code>brotli</code> , <code>zstd</code> , <code>bzip2</code> , <code>lz4</code> , <code>xz</code> , <code>snap</code> .
<code>compression::decompress(algo, data)</code>	Reverse.
<code>compression::decompress_auto(data)</code>	Auto-detect from magic bytes.
<code>compression::list()</code>	Supported algos.

Archive

Function	Description
<code>archive::create(dest_path, sources_array)</code>	Infer format from extension.
<code>archive::extract(src_path, dest_dir)</code>	Extract zip/tar/tar.gz/tar.xz/tar.bz2.

Examples

```
// Round-trip
let body = http("https://example.com/page.html").body;
let gz = compression::compress("gzip", body, 6);
print(`${body.len()} → ${gz.len()} bytes`);
let back = compression::decompress("gzip", gz);
assert(back.len() == body.len(), "round-trip");

// Zip up
archive::create("/tmp/backup.zip", ["file1.txt", "dir1/"]);
archive::extract("/tmp/backup.zip", "/tmp/restore/");
```

Encryption bindings

age + PGP shell-out.

Function	Description
<code>encrypt::keygen()</code>	New X25519 age keypair. Returns <code>#{ public, private }</code> .
<code>encrypt::encrypt(data, recipients)</code>	<code>recipients</code> = array of age-pub strings.
<code>encrypt::decrypt(data, identity)</code>	<code>identity</code> = secret key string.
<code>encrypt::encrypt_passphrase(data, passphrase)</code>	Script-based.
<code>encrypt::decrypt_passphrase(data, passphrase)</code>	—

Examples


```

let kp = encrypt::keygen();
print(`public: ${kp.public}`);
file_write_all("~/age/identity", kp.private);

let plaintext = "highly classified";
let ct = encrypt::encrypt(plaintext, [kp.public]);
let pt = encrypt::decrypt(ct, kp.private);
assert(pt.to_string() == plaintext, "round-trip");

let ct2 = encrypt::encrypt_passphrase("hello", "correct horse battery
staple");

```

Check-digit binding

Function	Description
<code>checkdigit::verify(algo, input)</code>	true / false.
<code>checkdigit::create(algo, partial)</code>	Computed full value.
<code>checkdigit::list()</code>	All 60+ algorithm names.

Examples

```

assert(checkdigit::verify("isbn13", "9780131103627"), "valid ISBN");

let full = checkdigit::create("ean13", "400638133393");
print(full);           // 4006381333931

for algo in checkdigit::list() {
    print(algo);
}

```

Sample-data binding

Function	Description
<code>sample::get(name)</code>	String content of a named sample.
<code>sample::list()</code>	Array of sample names.

Examples

```
let u = sample::get("json-user");
print(u);

for s in sample::list() { print(s); }
```

JWT binding

Function	Description
<code>jwt::view(token)</code>	Decode header + claims (no signature check).
<code>jwt::sign(payload_map, opts)</code>	opts: alg, secret or key path.
<code>jwt::validate(token, opts)</code>	opts: alg, secret or key / pubkey.

Examples

```
let t = jwt::sign({ sub: "alice", exp: now() + 3600 }, {
  alg: "HS256",
  secret: env("JWT_SECRET"),
});

let ok = jwt::validate(t, { alg: "HS256", secret: env("JWT_SECRET") });
assert(ok, "validation failed");

let parts = jwt::view(t);
print(`alg: ${parts.header.alg}, sub: ${parts.payload.sub}`);
```

Email-protection binding

Function	Description
<code>email::spf(domain)</code>	SPF record + validation.
<code>email::dmarc(domain)</code>	DMARC policy.
<code>email::dkim(domain, selector)</code>	DKIM record for selector.
<code>email::mta_sts(domain)</code>	MTA-STS policy.
<code>email::bimi(domain [, selector])</code>	BIMI record.
<code>email::tls_rpt(domain)</code>	SMTP TLS-RPT record.

Examples

```

let s = email::spf("example.com");
if s.valid { print("SPF OK"); } else { print(`SPF: ${s.error}`); }

let d = email::dmarc("example.com");
print(`policy: ${d.policy}`);

let dk = email::dkim("example.com", "selector1");
print(dk.record);

```

Netstatus binding

Function	Description
<code>netstatus::run()</code>	Execute the probe bundle defined in <code>~/.recon/config.toml</code> (layered — see Configuration files) [<code>netstatus</code>].

```

let r = netstatus::run();
for p in r.probes {
    print(`${p.name}: ${p.status}`);
}

```

Text-encoding binding

Function	Description
<code>text::decode(blob, charset)</code>	Decode bytes → string.
<code>text::encode(string, charset)</code>	Encode string → bytes.
<code>text::detect(blob)</code>	Auto-detect (<code>chardetng</code>). Returns charset label.
<code>text::normalize_newlines(s, style)</code>	style = <code>unix</code> / <code>windows</code> / <code>mac</code> .
<code>text::list_charsets()</code>	Supported labels.

Examples

```

let bytes = file_read("greek.txt");
let charset = text::detect(bytes);
let utf8 = text::decode(bytes, charset);
file_write_all("greek-utf8.txt", text::encode(utf8, "utf-8"));

let crlf = text::normalize_newlines("line1\nline2\n", "windows");

```

String helpers

PHP-style free functions for string manipulation. All are top-level callables — `trim(s)` reads the same as in PHP — and co-exist with Rhai's existing String methods (which keep working).

Function	Description
<code>trim(s) / trim(s, mask)</code>	Strip whitespace, or any character in <code>mask</code> , from both ends.
<code>ltrim(s) / ltrim(s, mask)</code>	Strip from the left end only.
<code>rtrim(s) / rtrim(s, mask)</code>	Strip from the right end only.
<code>strrev(s)</code>	Reverse a string by Unicode codepoints (accented letters and emoji stay intact).
<code>strip_html(s)</code>	Remove every <code><...></code> segment. Quoted attribute values are respected; HTML entities pass through.
<code>nl2br(s)</code>	Insert <code>
</code> before each <code>\n</code> , <code>\r\n</code> , or <code>\r</code> . HTML5 form (no trailing slash). Preserves the original newline.
<code>br2nl(s)</code>	Replace <code>
</code> / <code>
</code> / <code>
</code> (any case, any inner whitespace) with <code>\n</code> . If the tag is immediately followed by an EOL, that EOL is kept — so <code>nl2br</code> ↔ <code>br2nl</code> round-trips.
<code>preg_match(pattern, subject)</code>	Returns an Array of capture strings: index 0 is the whole match, 1+ are groups. Empty array if no match.
<code>preg_replace(pattern, replacement, subject)</code>	Replace every match. <code>\$1</code> / <code>\${name}</code> in <code>replacement</code> expand to captures.
<code>arr.join(sep) / join(arr, sep)</code>	Concatenate an Array's elements with <code>sep</code> between them. Non-string elements are stringified via <code>to_string</code> .
<code>sprintf(fmt) / sprintf(fmt, arg) / sprintf(fmt, [a, b, ...])</code>	Format and return a String.
<code>printf(fmt) / printf(fmt, arg) / printf(fmt, [a, b, ...])</code>	Format and write to stdout. Returns the byte count.
<code>urlencode(s) / urldecode(s)</code>	RFC 3986 percent-encoding for query params and form values. <code>urldecode</code> errors on malformed <code>%xx</code> sequences.

Function	Description
<code>base64_encode(s) / base64_encode(blob)</code>	Standard base64 with <code>=</code> padding. String input is encoded as UTF-8 bytes.
<code>base64_decode(s)</code>	Decode standard base64 to a Blob. Convert with <code>text::decode(b, "utf-8")</code> if you want a String.
<code>html_entity_decode(s)</code>	Decode HTML entities (<code>&amp;</code> , <code>&lt;</code> , <code>&#x27;</code> , numeric refs). Natural follow-up after <code>strip_html</code> .
<code>str_pad(s, width) / str_pad(s, width, pad) / str_pad(s, width, pad, side)</code>	Pad to <code>width</code> characters with <code>pad</code> (default space). <code>side</code> is <code>"left"</code> , <code>"right"</code> (default), or <code>"both"</code> . <code>Width ≤ length</code> leaves the string alone.
<code>lpad(s, width [, pad]) / rpad(s, width [, pad])</code>	Bare-name aliases for left/right pad with space (or <code>pad</code>) as the fill.
<code>dirname(path)</code>	POSIX <code>dirname</code> . Strips trailing slashes, returns everything before the last <code>/</code> . Returns <code>."</code> for a bare filename, <code>"/</code> for a rooted single-component path.
<code>basename(path) / basename(path, suffix)</code>	POSIX <code>basename</code> . Optional <code>suffix</code> is trimmed from the result (only when it doesn't equal the whole name).
<code>date_format(ts, fmt) / date_format(ts, fmt, tz)</code>	Format a Unix timestamp via chrono's <code>strftime</code> spec. <code>tz</code> defaults to UTC; pass <code>"local"</code> for the system timezone.

Regex patterns accept either raw form (`"foo\\d+"`) or PHP-style delimited form (`"/foo\\d+/i"`) with the `i` / `m` / `s` / `x` flags.

`printf` / `sprintf` specifiers: `d i u o x X b f e E g G s c %`. Flags: `-` (left-align), `0` (zero-pad), `+` (force sign), space (space-sign), `#` (alt form for `o` / `x` / `X` / `b`). Width and precision are both supported. Rhai has no variadic concept; pass `[a, b, c]` for multi-arg formats.

Examples

```

print(trim("  hello  "));          // "hello"
print(ltrim("...path", "."));      // "path"
print(rtrim("file.log", ".log"));  // "file"

print(strrev("café"));              // "éfac"
print(strip_html("<p>plain <b>text</b></p>")); // "plain text"
print(nl2br("a\nb"));              // "a<br>\nb"

let caps = preg_match("/^Host:\\s*(.+)$/i", "Host: example.com");
print(caps);                      // ["Host: example.com",
"example.com"]
print(preg_replace("\\s+", "-", "a b c")); // "a-b-c"

print(["a", "b", "c"].join(", ")); // "a, b, c"
print(sprintf("%-10s %5d", ["alpha", 42])); // "alpha          42"
print(sprintf("hex=%#x bin=%08b", [255, 10])); // "hex=0xff
bin=00001010"
printf("pi=%.4f\n", 3.14159265);    // writes "pi=3.1416\n"

print(urlencode("hello world & friends?")); //
"hello%20world%20%26%20friends%3F"
print(base64_encode("hello"));       // "aGVsbG8="
print(text::decode(base64_decode("aGVsbG8="), "utf-8")); // "hello"
print(html_entity_decode("&lt;b&gt;A &amp; B&lt;/b&gt;")); // "<b>A &
B</b>"

print(str_pad("42", 6, "0", "left")); // "000042"
print(rpad("hi", 5, "."));           // "hi..."
print(dirname("/var/log/recon.log")); // "/var/log"
print(basename("/var/log/recon.log", ".log")); // "recon"

print(date_format(1700000000, "%Y-%m-%dT%H:%M:%SZ")); // "2023-11-
14T22:13:20Z"
print(date_format(now_ms() / 1000, "%a %d %b %Y", "local"));

```

jq filter

Apply jq-style filters to any Rhai Map / Array. Backed by the [jq](#) crate — full jq grammar including pipes, `select(...)`, `map(...)`, alternative `//`, arithmetic, and the standard-library functions.

Function	Description
<code>obj.jq(filter) / jq(obj, filter)</code>	First result of the filter, or <code>()</code> (Rhai unit) if no result.

Function	Description
<code>obj.jq_all(filter) / jq_all(obj, filter)</code>	Every result as an Array. Empty Array if nothing matches.

The split avoids the ambiguity of a single auto-shaping method: a filter that "usually returns one match" wouldn't silently flip to an Array on the day it finds two.

Strings are NOT auto-parsed — chain `json_parse(s).jq(filter)` to start from JSON text. Filter parse errors and runtime errors both throw and are catchable with `try / catch`.

Example

```
let prs = [
  #{ number: 1, title: "Add foo", state: "open",  author: "alice" },
  #{ number: 2, title: "Fix bar", state: "closed", author: "bob"   },
  #{ number: 3, title: "Tweak",   state: "open",  author: "alice" },
];

// First / all
prs.jq(".[0].title");           // "Add
foo"
prs.jq_all(".[] | select(.state == \"open\") | .number"); // [1, 3]

// From raw JSON text
let raw = `{"items":[{"k":"a"}, {"k":"b"}]}`;
json_parse(raw).jq_all(".items[].k"); // ["a",
"b"]

// Catch a bad filter
try {
  prs.jq("bad syntax (");
} catch (e) {
  print(`caught: ${e}`);
}
```

git wrapper

First-class methods over the `git` CLI. Each method picks the right `--porcelain / --format` flag internally and parses the output into Rhai data. The `.run()` / `.run_text()` / `.run_json()` escape hatches expose anything not promoted to a method.

Function	Description
<code>git()</code> / <code>git(path)</code>	Construct a <code>Git</code> handle bound to the current working directory or an explicit repo path.

Function	Description
<code>g.status()</code>	Map { branch, upstream, ahead, behind, clean, staged, unstaged, untracked } from git status --porcelain=v2 --branch. staged / unstaged are arrays of { path, x_y } (renames also include original).
<code>g.is_clean()</code>	Convenience for <code>g.status().clean</code> .
<code>g.log(n) /</code> <code>g.log_range(rev_range)</code>	Array<Map { hash, short_hash, author, email, date, subject, body }>. ISO 8601 dates.
<code>g.diff()</code> / <code>g.diff(rev)</code>	Patch as a String.
<code>g.diff_stat()</code> / <code>g.diff_stat(rev)</code>	Map { files, insertions, deletions, per_file: [...] }.
<code>g.branch()</code>	Map { current, upstream, all }. upstream is () when no upstream is set.
<code>g.rev_parse(name)</code>	Resolve a ref to its full 40-char SHA.
<code>g.remote()</code>	Map of remote name → URL.
<code>g.add(path) /</code> <code>g.add([paths])</code>	Stage a path or an array of paths. Returns ().
<code>g.commit(message)</code>	Commit staged changes; returns { hash, short_hash, subject }. Throws on empty index or pre-commit hook failure.
<code>g.push()</code> / <code>g.push(remote) /</code> <code>g.push(remote, branch)</code>	Push. Returns ().
<code>g.pull()</code> / <code>g.pull(remote,</code> <code>branch)</code>	Pull. Returns ().
<code>g.checkout(name)</code>	Switch to a branch or commit. Returns ().
<code>g.run(args)</code>	Run any git args; sniffs JSON vs text.
<code>g.run_text(args) /</code> <code>g.run_json(args)</code>	Explicit text/JSON return forms.

Example


```

let g = git();

// Inspection.
print(g.branch().current);
print(g.is_clean());
for c in g.log(5) { print(`${c.short_hash} ${c.subject}`); }

// Mutation.
g.add("src/foo.rs");
let c = g.commit("fix: tighten the foo path");
print(`new commit ${c.short_hash}`);
g.push();

// Escape hatch.
let log = g.run_text("log --oneline -3");

```

Errors throw on non-zero exit; scripts use `try / catch` to recover. Composes on top of `std::process::Command` directly rather than going through the `shell()` binding.

gh wrapper

First-class methods over the GitHub CLI. Each method picks the right `--json <fields>` flag and parses the output into Rhai Maps and Arrays.

Auto-account-switch

The email-to-handle mapping is loaded from the `[gh.accounts]` table of the layered `config.toml` (see [Configuration files](#)). System and user layers are deep-merged with user winning per-entry. Removing the standalone `gh-accounts.toml` (shipped briefly in 0.89.0) was a breaking change in 0.90.0; any 0.89.0 users with a populated file need to copy the entries into `~/.recon/config.toml` under `[gh.accounts]`.

The switch is cached per `Gh` value (atomic check on every call, no redundant switches). `auth_status()` is the lone method that does NOT trigger auto-switch — useful when querying which account is active.

Methods

Function	Description
<code>gh()</code> / <code>gh("owner/name")</code>	Construct a <code>Gh</code> handle. The owner/name form adds <code>--repo <s></code> to every call.
<code>h.pr_list()</code> / <code>h.pr_list(opts)</code>	<code>Array<PR Map> . opts: state / author / label / limit</code> . <code>label</code> accepts string

Function	Description
	or array.
<code>h.pr_view(number)</code>	PR detail Map (body, labels, reviewDecision, mergeable, ...).
<code>h.pr_create(opts)</code>	Returns <code>{ number, url }</code> . <code>opts</code> : <code>title</code> (required), <code>body</code> OR <code>body_file</code> (mutex), <code>base</code> , <code>head</code> , <code>draft</code> , <code>reviewer</code> , <code>label</code> .
<code>h.pr_merge(number) /</code> <code>h.pr_merge(number, opts)</code>	<code>opts</code> : <code>method</code> (merge / squash / rebase), <code>delete_branch</code> , <code>auto</code> .
<code>h.pr_close(number) /</code> <code>h.pr_comment(number, body)</code>	Close or comment.
<code>h.issue_list() / h.issue_list(opts) /</code> <code>h.issue_view(number) /</code> <code>h.issue_create(opts) /</code> <code>h.issue_comment(number, body)</code>	Same shape as PR methods, scoped to issues. <code>issue_create</code> <code>opts</code> : <code>title</code> (required), <code>body</code> OR <code>body_file</code> , <code>label</code> , <code>assignee</code> .
<code>h.release_list() /</code> <code>h.release_view(tag)</code>	List or detail.
<code>h.release_create(tag, opts)</code>	Returns <code>{ url, tag }</code> . <code>opts</code> : <code>title</code> , <code>notes</code> OR <code>notes_file</code> (mutex), <code>generate_notes</code> , <code>draft</code> , <code>prerelease</code> , <code>target</code> .
<code>h.repo_view() / h.repo_view(spec)</code>	Repo metadata Map.
<code>h.run_list() / h.run_list(opts) /</code> <code>h.run_view(id)</code>	Workflow runs. <code>run_list</code> <code>opts</code> : <code>workflow</code> , <code>branch</code> , <code>status</code> , <code>limit</code> .
<code>h.auth_status()</code>	Map <code>{ host, account, scopes }</code> . Does NOT trigger auto-switch.
<code>h.run(args) / h.run_text(args) /</code> <code>h.run_json(args)</code>	Escape hatches, same shape as the git wrapper.

Example

```

let h = gh();

// Guard.
let auth = ();
try { auth = h.auth_status(); } catch (e) { print(`gh not configured:
${e}`); return; }
print(`active: ${auth.account}`);

// PR list + view.
for p in h.pr_list({ state: "open", limit: 5 }) {
    print(`#${p.number} ${p.title} by ${p.author.login}`);
}

// Create a release.
let r = h.release_create("v0.89.0", #{
    generate_notes: true,
    title: "0.89.0 - jq + git + gh",
});
print(`released at ${r.url}`);

// Catch "not found".
try {
    h.pr_view(999999);
} catch (e) {
    print(`pr not found: ${e}`);
}

```

Errors throw on non-zero exit with stderr truncated to ~2KB.

Whois binding

Function	Description
<code>whois(domain)</code>	Two-hop whois (registry then registrar). Returns raw text.

```

let w = whois("example.com");
print_raw(w);

```

IPFS binding

Function	Description
<code>ipfs(url)</code>	Fetch <code>ipfs://</code> or <code>ipns://</code> via the configured gateway.

```
let data = ipfs("ipfs://QmYwAPJzv5CZsnA625s3Xf2nemtYgPpHdWEz79ojWnPbdG");
print(data.body.len());
```

Agent-browser binding

Exposed as a static module. Requires `agent-browser` on `$PATH`.

Function	Description
<code>agentBrowser::available</code>	Bool.
<code>agentBrowser::version</code>	Version string.
<code>agentBrowser::open(url) / open(url, opts)</code>	Navigate to URL. <code>opts</code> overrides global options for this call (0.75.0).
<code>agentBrowser::close() / close_all()</code>	End the current session / close every session.
<code>agentBrowser::back() / forward() / reload()</code>	Browser history navigation.
<code>agentBrowser::click(selector) / dblclick(selector)</code>	Click / double-click. Selector is CSS, XPath, or <code>@ref</code> .
<code>agentBrowser::hover(selector) / focus(selector)</code>	Hover / focus an element.
<code>agentBrowser::check(selector) / uncheck(selector)</code>	Toggle a checkbox.
<code>agentBrowser::fill(selector, text)</code>	Clear and fill a form field.
<code>agentBrowser::type_text(selector, text)</code>	Type into a field (per-character events). Named <code>type_text</code> because <code>type</code> is reserved in Rhai.
<code>agentBrowser::keyboard_type(text) / keyboard_insert(text)</code>	Keyboard-level typing into the focused element.
<code>agentBrowser::press(key)</code>	Key press (e.g. "Enter", "Tab", "Control+a").
<code>agentBrowser::scroll(dir) / scroll(dir, px)</code>	Scroll the page.
<code>agentBrowser::scrollintoview(selector)</code>	Scroll an element into the viewport.
<code>agentBrowser::wait(selector_or_ms)</code>	Wait for element or a millisecond duration (string or int).

Function	Description
<code>agentBrowser::screenshot()</code> / <code>screenshot(path)</code> / <code>screenshot(path, opts)</code>	Save PNG (default path or explicit).
<code>agentBrowser::pdf(path)</code> / <code>pdf(path, opts)</code>	Save PDF.
<code>agentBrowser::snapshot()</code> / <code>snapshot(interactive_bool)</code> / <code>snapshot(opts)</code> / <code>snapshot(interactive, opts)</code>	Accessibility-tree dump. <code>interactive=true</code> filters to interactive elements only.
<code>agentBrowser::eval_js(js)</code> / <code>eval_js(js, opts)</code>	Execute JS, return the parsed value. Named <code>eval_js</code> because <code>eval</code> is reserved in Rhai.
<code>agentBrowser::get(what)</code> / <code>get(what, selector)</code>	Read page or element data. <code>what</code> ∈ <code>text</code> / <code>html</code> / <code>value</code> / <code>attr</code> <code><name></code> / <code>title</code> / <code>url</code> / <code>count</code> / <code>box</code> / <code>styles</code> / <code>cdp-url</code> . Returns a JSON object whose key matches <code>what</code> .
<code>agentBrowser::find(locator, value, action)</code> / <code>find(locator, value, action, text)</code>	Semantic-locator find + act. <code>locator</code> ∈ <code>role</code> / <code>text</code> / <code>label</code> / <code>placeholder</code> / <code>alt</code> / <code>title</code> / <code>testid</code> / <code>first</code> / <code>last</code> / <code>nth</code> . <code>action</code> ∈ <code>click</code> / <code>dblclick</code> / <code>hover</code> / <code>focus</code> / <code>fill</code> / <code>type</code> / <code>check</code> / <code>unchecked</code> . The 4-arg form passes <code>text</code> for fill/type.
<code>agentBrowser::is_visible(selector)</code> / <code>is_enabled(selector)</code> / <code>is_checked(selector)</code>	Predicate checks. Returns <code>bool</code> . Errors if no element matches — use <code>get("count", sel)</code> for an existence check that doesn't raise.
<code>agentBrowser::cmd(args_array)</code>	Raw passthrough to the CLI. Escape hatch for any subcommand without a typed wrapper (cookies, storage, tabs, network, mouse, etc.).

Existence check

To test whether a selector matches anything without raising an error:

```
fn exists(sel) {
    agentBrowser::get("count", sel).count > 0
}

if exists("#login-form") {
    agentBrowser::fill("#user", "alice");
}
```

`is_visible` / `is_enabled` / `is_checked` raise a Rhai error when no element matches, which aborts the script unless wrapped in `try { ... } catch { ... }`. `get("count", sel)` always returns `{ count: N }`, so `count == 0` is the no-match signal.

Examples

```
if !agentBrowser::available { return 2; }

agentBrowser::open("https://example.com/login");
agentBrowser::fill("#user", "alice");
agentBrowser::fill("#pass", "s3cr3t");
agentBrowser::click("button[type=submit]");
agentBrowser::screenshot("/tmp/after-login.png");
agentBrowser::close();
```

See `script/agent-browser-find.rhai`, `script/agent-browser-interaction.rhai`, `script/agent-browser-inspect.rhai`, `script/agent-browser-navigation.rhai`, `script/agent-browser-pdf.rhai`, and `script/agent-browser-cmd.rhai` for focused demos of each part of the surface.

Global options (0.75.0)

`agent-browser` accepts ~25 global launch / security / session options (see `agent-browser --help`). All are exposed as `opts-map` keys via `agentBrowser::set_default_options(opts)` (module-level defaults that apply to every binding call) plus per-call overrides on launch verbs.

Rhai key	agent-browser flag	Type
<code>ignore_https_errors</code>	<code>--ignore-https-errors</code>	bool
<code>allow_file_access</code>	<code>--allow-file-access</code>	bool
<code>headed</code>	<code>--headed</code>	bool
<code>auto_connect</code>	<code>--auto-connect</code>	bool
<code>annotate</code>	<code>--annotate</code>	bool
<code>no_auto_dialog</code>	<code>--no-auto-dialog</code>	bool
<code>content_boundaries</code>	<code>--content-boundaries</code>	bool

Rhai key	agent-browser flag	Type
confirm_interactive	--confirm-interactive	bool
verbose / quiet / debug / json	--verbose / --quiet / --debug / --json	bool
session	--session	string
session_name	--session-name	string
executable_path	--executable-path	string
user_agent	--user-agent	string
proxy	--proxy	string
proxy_bypass	--proxy-bypass	string
state	--state	string
profile	--profile	string
provider	--provider	string
device	--device	string
color_scheme	--color-scheme	string
engine	--engine	string
model	--model	string
config	--config	string
screenshot_dir	--screenshot-dir	string
screenshot_format	--screenshot-format	string
download_path	--download-path	string
allowed_domains	--allowed-domains	string
action_policy	--action-policy	string
confirm_actions	--confirm-actions	string
cdp	--cdp	int
screenshot_quality	--screenshot-quality	int
max_output	--max-output	int
extension	--extension (repeatable)	string or array
browser_args	--args (repeatable)	string or array

Rhai key	agent-browser flag	Type
headers	--headers <json>	string or map

Module functions:

- `agentBrowser::set_default_options(opts: Map)` — store the defaults.
- `agentBrowser::clear_default_options()` — reset to empty.
- `agentBrowser::default_options() → Map` — read current defaults.

Per-call opts are accepted on `open`, `screenshot`, `snapshot`, `pdf`, and `eval`. They concatenate after defaults; agent-browser's flag parser uses last-wins for repeated single-value flags, so per-call options override defaults.

```
agentBrowser::set_default_options({
  ignore_https_errors: true,
  user_agent: "Recon/0.75",
  proxy: "http://proxy:3128",
});

agentBrowser::open("https://self-signed.example");
agentBrowser::click("#login");

// Per-call override:
agentBrowser::open("https://other.example", #{ user_agent: "Other/1.0" });

// Headers as a Rhai map (auto-serialized to JSON):
agentBrowser::open("https://api.example", #{
  headers: #{ Authorization: "Bearer x" },
});
```

SQLite binding

Function	Description
<code>sqlite(spec) / sqlite(spec, mode)</code>	Open. spec = path, <code>":memory:"</code> , or an alias. mode = <code>ro</code> <code>rw</code> (default) <code>rwc</code> .
<code>.query(sql) / .query(sql, params)</code>	Array of Maps.
<code>.query_one(sql, params)</code>	First row as a Map (or <code>()</code>).
<code>.query_value(sql, params)</code>	Single scalar value.
<code>.exec(sql, params)</code>	Row count affected.

Examples

```
let db = sqlite(":memory:");
db.exec("CREATE TABLE users (id INTEGER PRIMARY KEY, name TEXT);");
db.exec("INSERT INTO users (name) VALUES (?)", ["alice"]);
db.exec("INSERT INTO users (name) VALUES (?)", ["bob"]);

let rows = db.query("SELECT id, name FROM users ORDER BY id");
for r in rows {
  print(`${r.id}: ${r.name}`);
}

let count = db.query_value("SELECT COUNT(*) FROM users");
print(`total: ${count}`);
```

Document-conversion bindings

0.58.0. comrak for HTML; agent-browser for PDF.

Function	Description
<code>md_to_html(md)</code> / <code>md_to_html(md, opts)</code>	Markdown (string or Blob) → HTML string.
<code>md_to_pdf(md, dest)</code> / <code>md_to_pdf(md, dest, opts)</code>	Markdown → PDF at <code>dest</code> . Needs agent-browser.
<code>html_to_pdf(html, dest)</code>	HTML → PDF at <code>dest</code> . Needs agent-browser.
<code>html_to_text(html)</code> / <code>html_to_text(html, # {width, color, no_links})</code>	Render an HTML string to text. <code>color</code> is "auto" / "always" / "never" (default "never" in scripts). <code>no_links</code> : true drops the link footnote list (plain) / inline link styling (coloured).

Opts map

Key	Default	Description
<code>toc</code>	false	Inject linkable TOC.
<code>toc_depth</code>	3	Include H1..H N.
<code>toc_title</code>	"Contents"	—
<code>title</code>	document default	<title> + PDF metadata.
<code>css</code>	—	Additional CSS (appended after default).
<code>no_default_css</code>	false	Skip bundled default CSS.

Key	Default	Description
<code>gfm</code>	false	GitHub-flavored extensions.

Examples

```
let md = file_read("README.md").to_string();

// md → html
let html = md_to_html(md, #{ toc: true, toc_depth: 3, gfm: true, title:
"README" });
file_write_all("/tmp/readme.html", html);

// md → pdf
md_to_pdf(md, "/tmp/readme.pdf", #{ toc: true, gfm: true, title: "README"
});

// html → pdf
let html = http("https://example.com/report.html").body.to_string();
html_to_pdf(html, "/tmp/report.pdf");

// Fully custom styling
md_to_pdf(md, "/tmp/styled.pdf", #{
  no_default_css: true,
  css: file_read("print.css").to_string(),
  toc: true,
  toc_title: "Table of Contents",
});

// html → text (lynx/w3m-style)
let text = html_to_text(page_html, #{ width: 100 });
let resp = http("https://example.com", #{ render: true }); // resp.body is
text
```

pdf_export_page

Signature	Returns	Description
<code>pdf_export_page(pdf, page)</code>	<code>Blob</code>	Render page as PNG bytes.
<code>pdf_export_page(pdf, page, opts)</code>	<code>Blob</code>	Render page; opts map drives format / viewport / scale / quality.
<code>pdf_export_page(pdf, page, dest)</code>	<code>()</code>	Write to dest path; format from extension.
<code>pdf_export_page(pdf, page, dest, opts)</code>	<code>()</code>	Write to dest with full options.

Opts map keys:

Key	Type	Default	Notes
viewport	string "WxH"	"1024x1366"	Chrome CSS-pixel viewport.
scale	int	2	Device scale factor.
quality	int 0–100	90	JPEG/WEBP quality.
format	"png" "jpeg" "webp"	inferred from dest extension, else png	Explicit override.

Examples:

```
// Write PNG, defaults
pdf_export_page("report.pdf", 1, "/tmp/cover.png");

// Larger WEBP
pdf_export_page("report.pdf", 1, "/tmp/cover.webp", #{
  viewport: "1920x2715",
  scale: 2,
  quality: 80,
});

// Return bytes
let png_bytes = pdf_export_page("report.pdf", 1);
let jpeg_bytes = pdf_export_page("report.pdf", 1, #{ format: "jpeg",
  quality: 70 });
```

Requires `pdftoppm` (poppler-utils) on PATH. Install via `brew install poppler` (macOS) or `apt install poppler-utils` (Debian/Ubuntu).

AI bindings (`ai::*`)

The `ai::*` namespace dispatches a prompt to one of several subprocess-driven agent CLIs. Built-in backends: `claude` (Anthropic Claude Code), `codex` (OpenAI Codex), `copilot` (GitHub Copilot CLI), `gemini` (Google Gemini CLI). A user-defined `cmd` backend covers anything else. An HTTP backend (Anthropic Messages, OpenAI Chat Completions) is reserved as a follow-up — the builder API is forward-compatible.

Function	Returns	Description
<code>ai::ask(prompt)</code>	string	One-shot: <code>request().prompt(p).send()</code> .
<code>ai::request()</code>	builder	Fresh <code>AiRequest</code> builder.

Builder methods on an `AiRequest` :

Method	Behaviour
<code>.backend(name)</code>	Select backend (<code>claude</code> / <code>codex</code> / <code>copilot</code> / <code>gemini</code> / config-defined).
<code>.model(name)</code>	Pass-through to the backend's <code>--model</code> flag or equivalent.
<code>.system(s)</code>	System prompt; singleton.
<code>.context(s)</code>	Append a context block; multiple calls accumulate.
<code>.prompt(s)</code> / <code>.user(s)</code>	Set the current user turn (singleton).
<code>.assistant(s)</code>	Append a prior assistant turn (for multi-turn replay).
<code>.max_tokens(n)</code> , <code>.temperature(f)</code>	Hints; backend honours when available.
<code>.timeout(secs)</code>	Wall-clock kill switch; default 60.
<code>.send()</code>	Returns <code>string</code> — model's reply. Throws on failure.
<code>.send_full()</code>	Returns map: <code>.text</code> <code>.backend</code> <code>.model</code> <code>.duration_ms</code> <code>.exit_code</code> .

Backend selection

Three-layer precedence per `.send()` :

1. Per-request `.backend(name)` / `.model(name)` / `.timeout(secs)` .
2. Env vars: `RECON_AI_BACKEND` , `RECON_AI_MODEL` , `RECON_AI_TIMEOUT` .
3. `~/.recon/config.toml` (layered — see [Configuration files](#)) `[ai]` section.

No PATH fallback. If nothing selects a backend the request errors.

Config file

```

[ai]
default_backend = "claude"
default_model    = "sonnet"
timeout_secs     = 60

[ai.backends.claude]
model = "claude-sonnet-4-5"

[ai.backends.my-llm]
cmd      = ["my-llm-cli", "--print"]
model_flag = "--model"
system_flag = "--system"

```

The `[ai.backends.<name>]` block is the escape hatch for new agent CLIs without recompiling — name it, give it argv, optionally name its model and system flags. Scripts then write `req.backend("my-llm").model("v1")`.

Example

```

let req = ai::request()
    .system("You are a concise TLS expert.")
    .context("Cert subject CN: example.com")
    .context("Issuer: Let's Encrypt R3")
    .prompt("Is this a typical commercial cert?")
    .timeout(30);
print(req.send());

```

Errors

All `.send()` failures throw a Rhai script error prefixed `ai: :`

Tag	When
<code>ai: backend not configured</code>	no env, no config, no <code>.backend()</code>
<code>ai: backend '<name>' not found</code>	unknown built-in / config entry
<code>ai: CLI exited with status <N>:\n<stderr tail></code>	non-zero exit
<code>ai: timed out after <N>s</code>	subprocess killed by timeout
<code>ai: spawn failed: <io error></code>	CLI not on PATH, etc.
<code>ai: empty response from backend</code>	exit-zero but no stdout

Tag	When
ai: no user prompt – call <code>.prompt()/.user()</code> before <code>.send()</code>	builder validation
ai: cannot append assistant turn – last turn is already assistant	builder validation

Telemetry (`-v` / `-vv`)

Silent by default. When the script engine runs at verbosity 1 (`-v`), each **successful** `.send()` logs one line to stderr:

```
* ai: backend=claude model=sonnet duration=2.3s exit=0 chars_in=842
chars_out=412
```

`chars_in` is the total payload (system + context + multi-turn preamble + prompt); `chars_out` is the reply length. At verbosity 2 (`-vv`) a second line is added with the preamble length and the first 80 characters of stdout (newlines collapsed to spaces). Failed sends emit their own `ai: error` instead — no telemetry line. The reply on stdout is never affected.

Part IV — Appendices

Exit codes

Curl-compatible exit codes for common cases:

Code	Meaning
0	Success.
1	Generic protocol error.
2	Source-load error (<code>--compare</code> , etc.).
3	Agent-browser / Chrome unavailable (PDF features).
6	DNS resolution failed.
7	Connection refused.
22	HTTP $\geq$ 400 with <code>-f</code> (or <code>--fail-with-body</code>).
28	Operation timeout (<code>--max-time</code>).
35	TLS handshake error.
47	Redirect limit exceeded.
52	Empty reply from server.
55	Send error.
56	Receive error.
60	Peer cert cannot be authenticated.

Script-mode exit codes:

- `return N` from the top-level script becomes the process exit code (mod 256).
- Unhandled script errors → exit 1 (or the last `ProtocolExitCode` tag).

Environment variables

Variable	Read by	Description
<code>HTTP_PROXY</code> / <code>http_proxy</code>	<code>--proxy</code>	Proxy for http:// URLs.

Variable	Read by	Description
HTTPS_PROXY / https_proxy	--proxy	Proxy for https:// URLs.
ALL_PROXY / all_proxy	--proxy	Fallback proxy.
NO_PROXY / no_proxy	--noproxy	Bypass list.
SSL_CERT_FILE	--cacert	Extra trust-root bundle.
CURL_CA_BUNDLE	--cacert	Same as SSL_CERT_FILE.
RECON_NO_PAGER	--help / --examples	Disable paging.
RECON_IPFS_GATEWAY	--ipfs-gateway	Default IPFS gateway.
RECON_CONFIG	config file	Override path to config.toml .
AGENT_BROWSER_JSON	—	When 1 , agent-browser returns JSON (set internally by recon).
EDITOR	--editor	Editor binary for response inspection.
HOME	various	Used to locate ~/.recon/ .
PAGER	--help / --examples	Override default less -FRX .

Configuration file

~/.recon/config.toml (layered — see [Configuration files](#)) — commented skeleton created by recon --init.

Structure


```

# Default flag overrides
[defaults]
connect_timeout = 30
max_time = 60
user_agent = "recon/0.58.1 (auto)"

# Netstatus probe bundle
[netstatus]
probes = [
    { name = "dns", target = "8.8.8.8" },
    { name = "http", url = "https://example.com/" },
    { name = "tls", target = "example.com:443" },
]

# Per-host SNI → cert mapping for --serve-tls
[[serve_sni]]
host = "a.example.com"
cert = "/etc/certs/a.pem"
key = "/etc/certs/a.key"

[[serve_sni]]
host = "b.example.com"
cert = "/etc/certs/b.pem"
key = "/etc/certs/b.key"

```

~/.recon/ layout

Created by `recon --init`:

```

~/.recon/
├─ config.toml           # Commented skeleton; overrideable via
RECON_CONFIG
├─ script/              # Bare-word scripts (`recon --script NAME`)
│   └─ *.rhai
├─ jars/                # Persistent cookie jars
│   └─ <session>.db
├─ sni/                 # Per-host TLS cert + key pairs for --serve-
tls
│   └─ a.example.com/cert.pem
│   └─ a.example.com/key.pem
└─ hsts.txt             # Optional HSTS cache (curl-compatible)

```

Glossary

- **Sticky session** — a `browser()` handle whose cookies + headers persist across calls. Pair with `use_persistent_session` for cross-invocation persistence.
 - **Script defaults** — the frozen snapshot of CLI flags (`-H`, `-k`, `--connect-timeout`, ...) that every script binding inherits when the caller doesn't override via an opts map.
 - **Bare-word script** — one that can be invoked by name because it lives in `~/.recon/script/`, e.g. `recon --script http` → `~/.recon/script/http.rhai`.
 - **Protocol exit code** — recon's internal scheme for passing an exit code through a panic-style error path from protocol bindings (so a failed SMTP probe can exit with curl's `CURLE_URL_MALFORMAT`, for instance).
 - **agent-browser** — external CLI that wraps Chrome DevTools Protocol. Used for `--browser-screenshot`, `--md-to-pdf`, `--html-to-pdf`, and the `agentBrowser::*` script bindings.
-

End of manual. For the curated-example companion, run `recon --examples`. For topic-specific deep dives, run `recon --help <topic>`.